

Memory Model -2- (mem_map)

📅 2016-03-03 (http://jake.dothome.co.kr/mem_map/) 👤 Moon Young-il

(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

<kernel v5.15>

Memory Model -2- (mem_map)

Learn how to manage page frame numbers (PFNs) and mem_map by physics model.

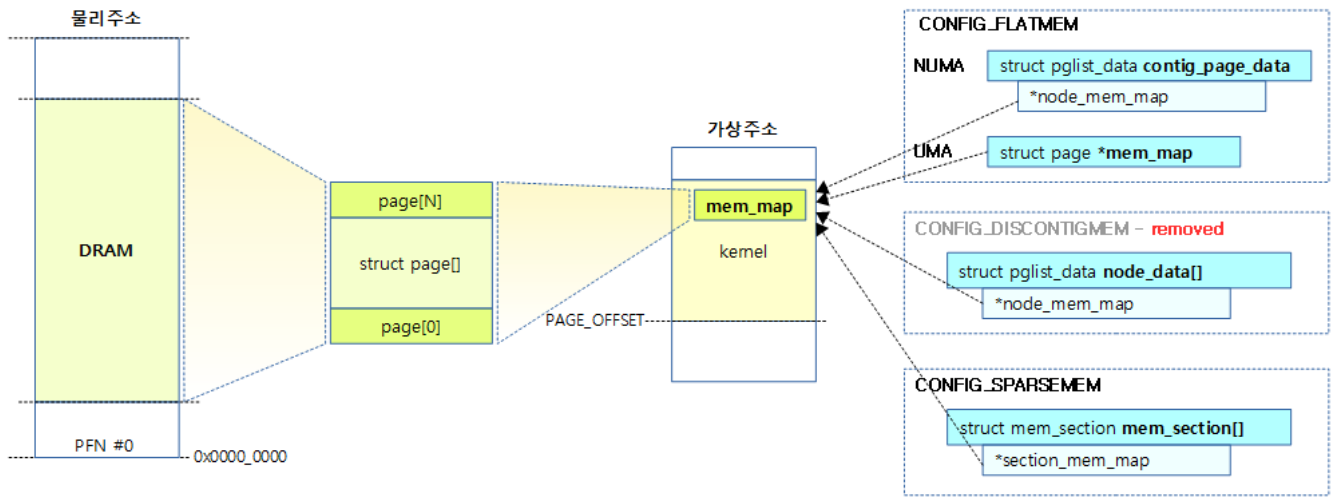
PFN(Page Frame Number)

A PFN is a page-by-page number used for a physical address, starting with zero. 0K is usually used for page units, and in the case of ARM4, there are systems that are set to 64K and 16K for the purpose of operating large-scale databases.

- e.g. 4K Page Standard: 0x0000_0000
 - The PFN value is 0.
- e.g. 4K Page Standard: 0x1234_5678
 - The PFN value is 0x12345.

mem_map

- mem_map is an array of page structures that contain information about all physical memory page frames.
 - In Linux, it started with an array of reference counters of type unsigned short, but gradually the number of members it needed increased, and it became the page structure of today.
- The approach is different for NUMA and UMA systems, and for each physical memory model.
- The path to the initialization function in mem_map varies by architecture and kernel configuration.
 - arm with flatmem
 - setup_arch() → paging_init() → bootmem_init() → zone_sizes_init() → free_area_init_node() (http://jake.dothome.co.kr/free_area_init_node)
 - arm/arm64 with sparsemem
 - setup_arch() → bootmem_init() → sparse_init()



(http://jake.dothome.co.kr/wp-content/uploads/2016/03/mem_map-1d.png)

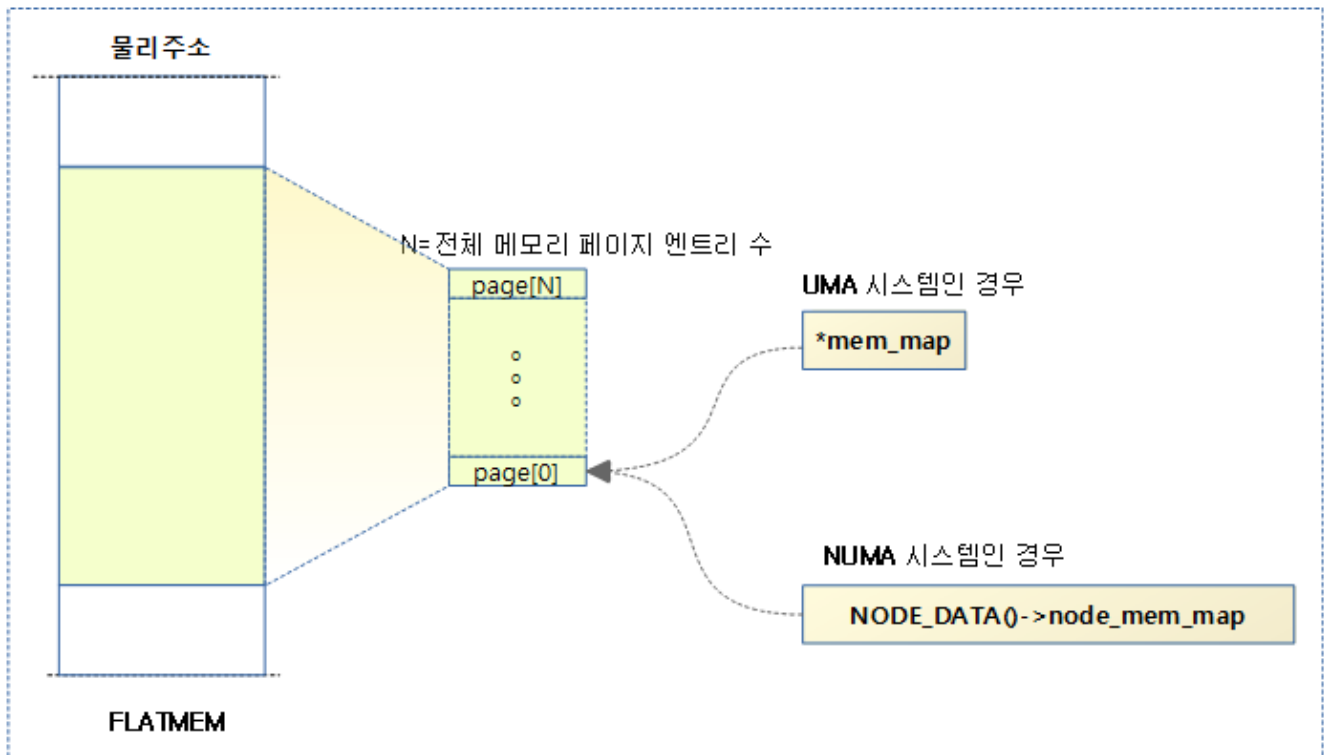
PFN vs page 변환 API

The Linux kernel uses very frequent conversions to access page structures with PFN values, and vice versa. The corresponding APIs are as follows.

- `pfn_to_page()`
 - Returns a pointer to the page structure with a pfn number.
- `page_to_pfn()`
 - Returns the pfn number as a pointer to the page structure.

Flat Memory with mem_map

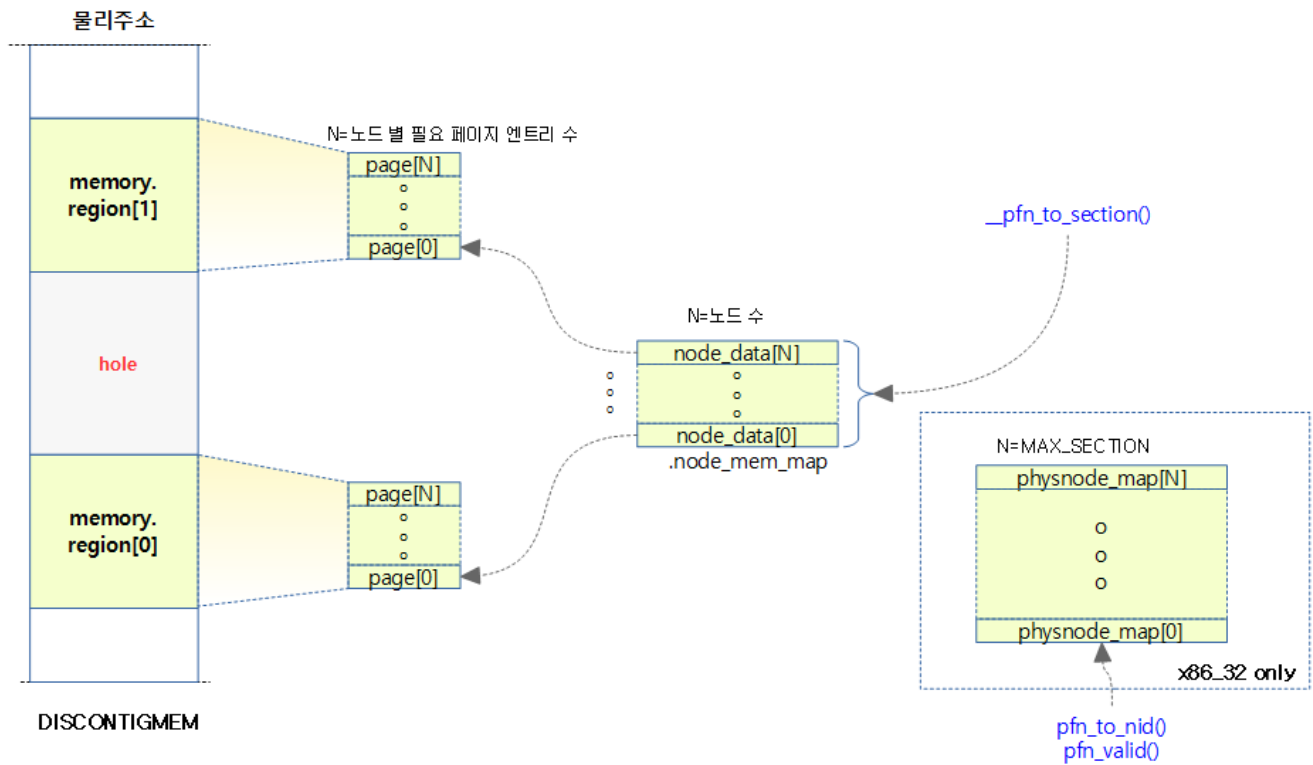
- When the `NEED_MULTIPLE_NODES` kernel option is not used, the `*mem_map` pointer variable points to an array of a single `page[]` structure.
- If you use the `FLAT_NODE_MEM_MAP` kernel option, `contig_page_data.node_mem_map` points to an array of `page[]` structures.



(<http://jake.dothome.co.kr/wp-content/uploads/2016/03/mm-8c.png>)

Discontiguous Memory with mem_map

- `node_data[].node_mem_map` refers to an array of `page[]` structures.
- In x86_32, a separate implementation of a table for section-to-node mapping is used.
- It only manages the memory it actually uses, and it doesn't manage the hole area at all.
- In the kernel mainline, this model has been completely removed so that it can no longer be used.
 - mm: remove CONFIG_DISCONTIGMEM
 (<https://github.com/torvalds/linux/commit/bb1c50d3967f69f413b333713c2718d48d1ab7ea#diff-6a0166b1d8bbb576287048e4de42eb7e8a1f118e357f407d3bdf7da9fc26d94d>)
 (2021, v5.14-rc1)



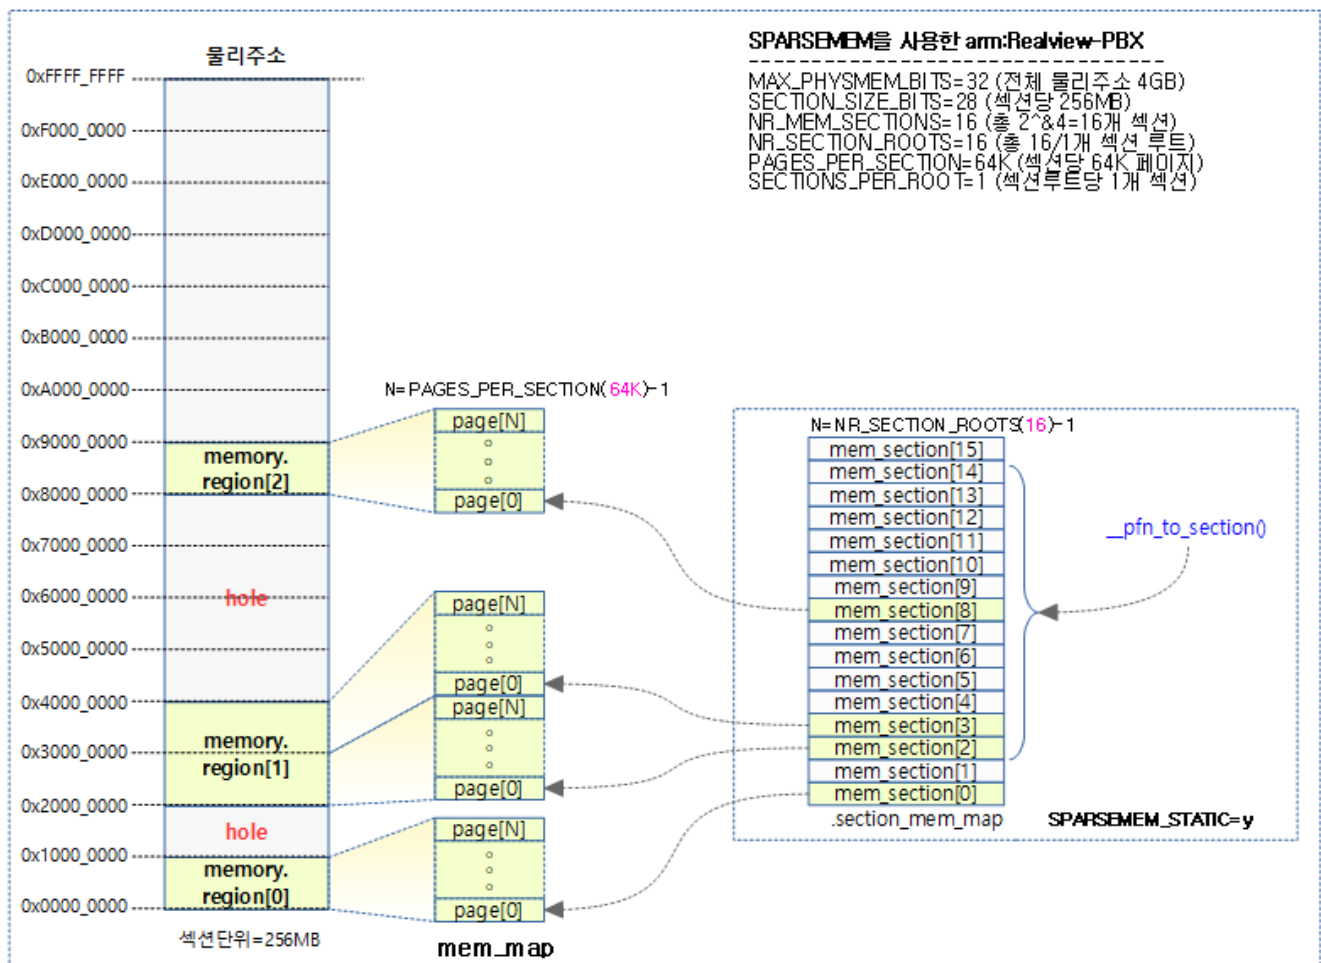
(<http://jake.dothome.co.kr/wp-content/uploads/2016/03/mm-6.png>)

Sparse Memory with mem_map

- It manages the `mem_map` through multiple sections, and uses `mem_section` structs to manage a single section.
- The size of a single section should be tens of MB ~ several GB in size.
 - Note: Not to be confused with the section terminology used in the page table.
 - In the case of the arm64 system, the initial 1G was used, but it was changed to a slightly smaller unit, the 128M.
 - Note: arm64/sparsemem: reduce SECTION_SIZE_BITS
(<https://github.com/torvalds/linux/commit/f0b13ee23241846f6f6cd0d119a8ac8059416ec4#diff-9590f22a80ad3a29d9c97bab548e481c48ac04e5ca5441c9360c69458829634c>) (2021, v5.12-rc1)
- A `mem_section` struct points to and manages a `PAGES_PER_SECTION` number of `page[]` arrays.
- There are two implementations of `mem_section` structs that manage sections to manage memory allocations.
 - `SPARSEMEM_STATIC`
 - On systems with small physical memory, such as 32-bit systems, `mem_section` structures are generated at compile time en masse.
 - `mem_section[][1]` is a double array, but does not use a second index.
 - `SPARSEMEM_EXTREME`
 - On systems with large physical memory sizes, such as 64-bit systems, `mem_section` structs are created and used at runtime when needed.
 - It is designed to reduce memory waste in case the hole size is very large.

- ****mem_section** Use a double pointer, and use it in the form of a `mem_section[SECTIONS_PER_ROOT][SECTIONS_PER_ROOT]`, and is used as follows:
 - For the first array index, it means the section root
 - `NR_SECTION_ROOTS` (number of section roots) = `NR_MEM_SECTIONS` (number of sections to use for total memory) / `SECTIONS_PER_ROOT` (number of `mem_section` structs that can fit in a 1 page frame)
 - The second array index refers to the `mem_section` index within the root section.
 - `SECTIONS_PER_ROOT` (the number of `mem_section` structs that can fit in a 1 page frame)

The following illustration shows a 32-bit arm – Realview-PBX board without using `SPARSEMEM_EXTREME`.

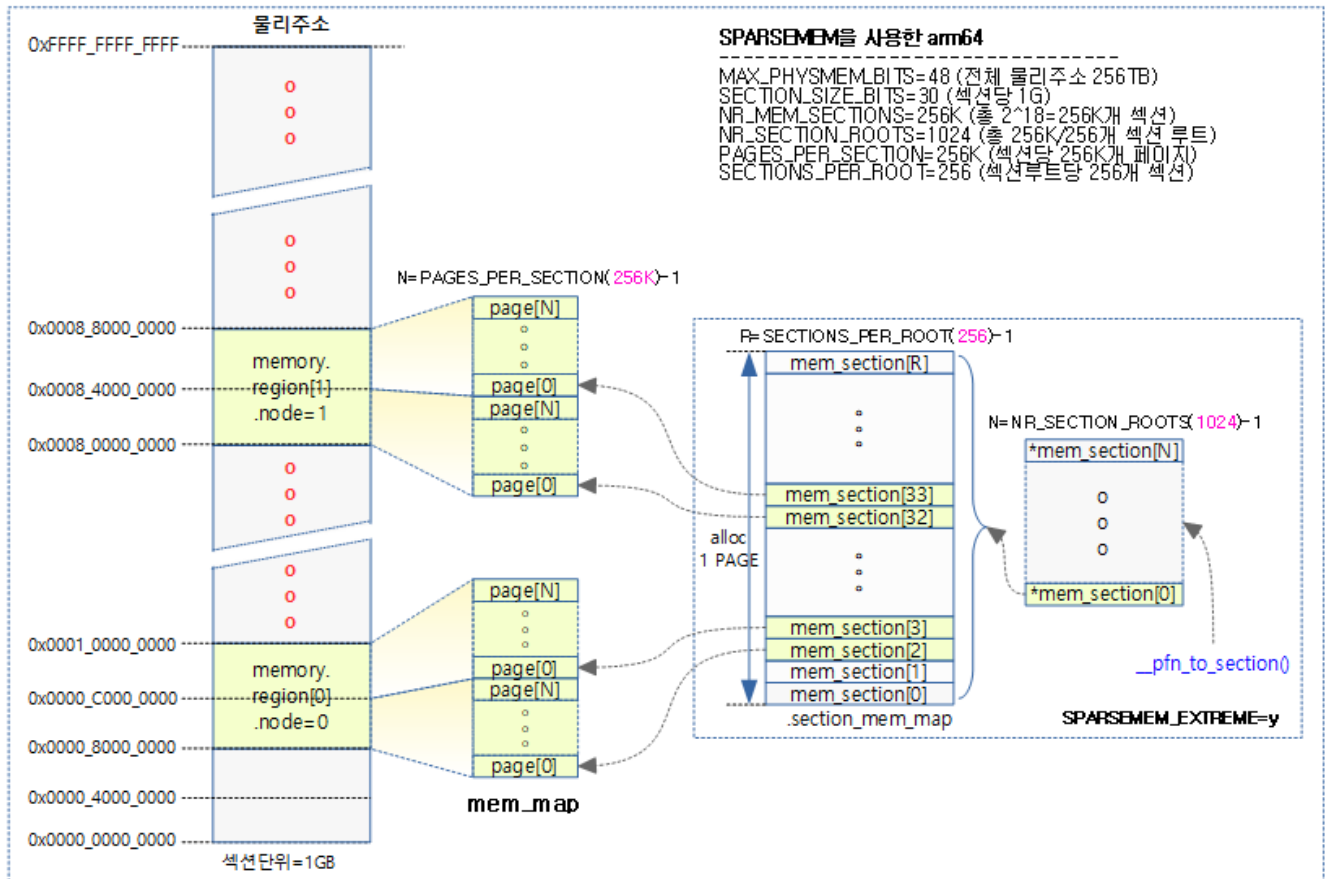


(<http://jake.dothome.co.kr/wp-content/uploads/2016/03/mm-7c.png>)

The following figure shows the use of 64G DRAM in the arm4 architecture, using `SPARSEMEM_EXTREME`.

- Below, we have shown an example of using 1G for the section size, but since it has recently been changed to 128M, this should be taken into account depending on the kernel version.
 - Note: arm64/sparsemem: reduce SECTION_SIZE_BITS
 (<https://github.com/torvalds/linux/commit/f0b13ee23241846f6f6cd0d119a8ac8059416ec4>)

#diff-9590f22a80ad3a29d9c97bab548e481c48ac04e5ca5441c9360c69458829634c) (2021, v5.12-rc1)



(<http://jake.dothome.co.kr/wp-content/uploads/2016/03/mm-9b.png>)

page descriptor

One page descriptor is assigned to every physical memory page. It is generated for every memory page, so it is size-sensitive. Therefore, in order to reduce the size as much as possible, it is designed to group the members managed in the page descriptor into a union type.

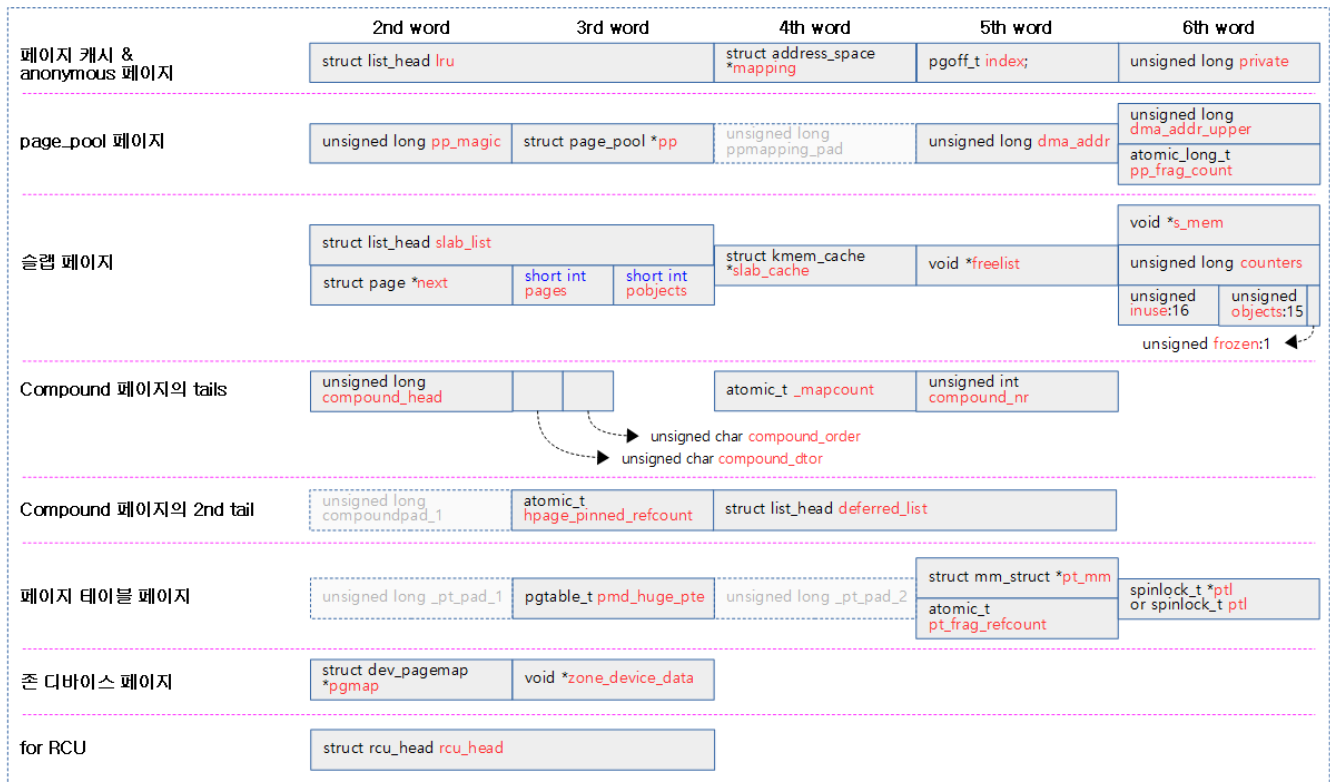
The following shows the page descriptor that works on a 32-bit system. (Use 32 bytes as the minimum configuration excluding all options.)

struct page 구조체 (32bits: word=4 bytes)

total: 32~44Bytes

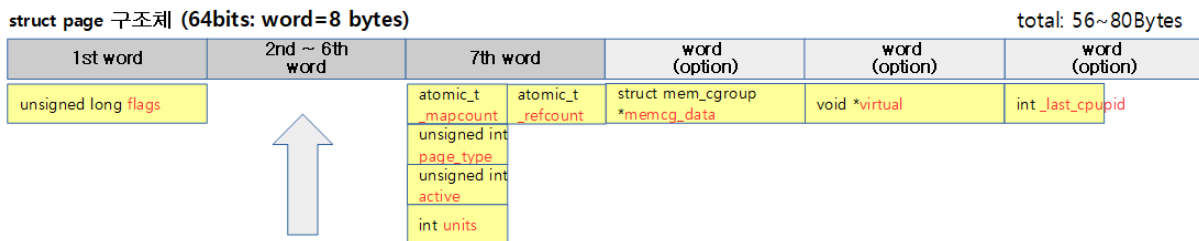
1st word	2nd ~ 6th word	7th word	8th word	word (option)	word (option)	word (option)
unsigned long flags	<div>↑</div>	atomic_t mapcount	atomic_t refcount	struct mem_cgroup *memcg_data	void *virtual	int last_cpupid
		unsigned int page_type				
		unsigned int active				
		int units				

(http://jake.dothome.co.kr/wp-content/uploads/2016/03/mem_map-3b.png)

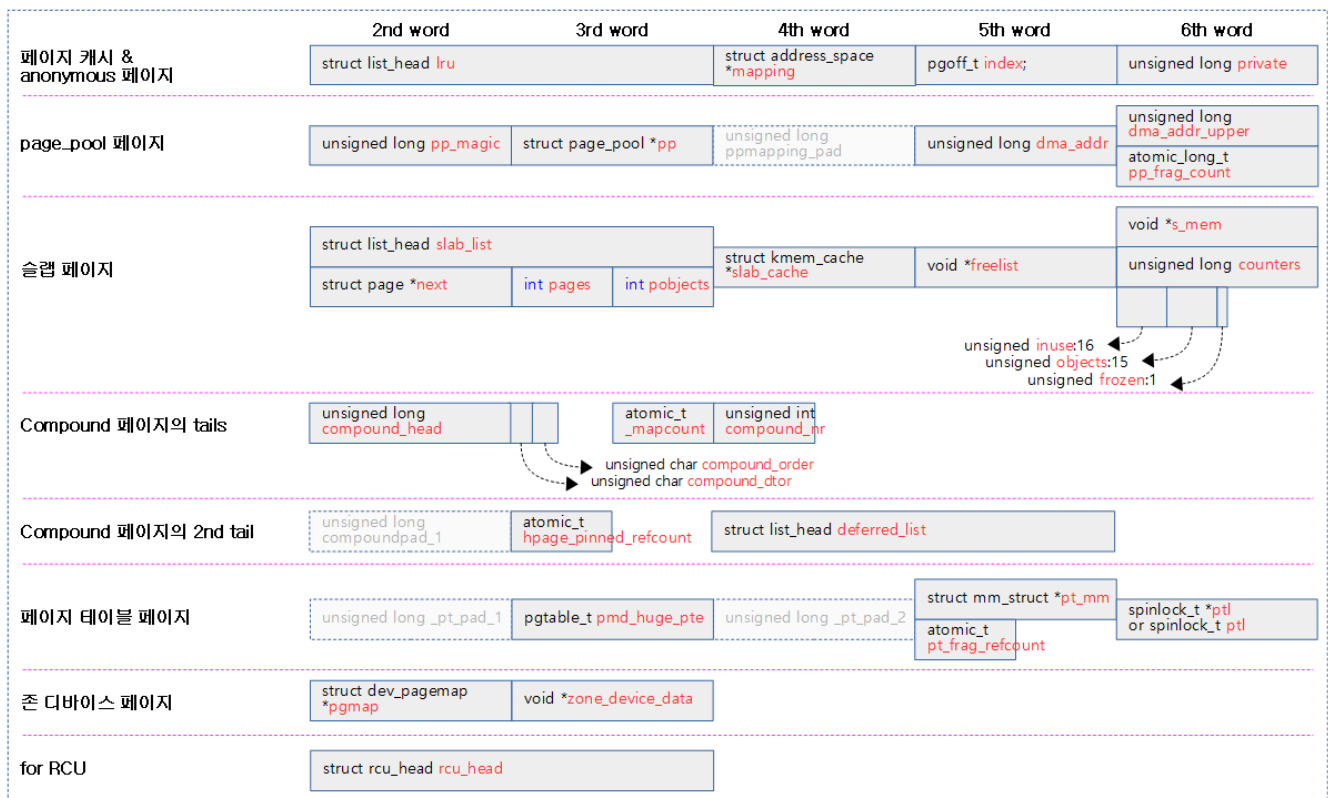


(http://jake.dothome.co.kr/wp-content/uploads/2016/03/mem_map-4b.png)

The following shows the page descriptor that works on a 64-bit system. (The default option is 64 bytes.)



(http://jake.dothome.co.kr/wp-content/uploads/2016/03/mem_map-5a.png)



(http://jake.dothome.co.kr/wp-content/uploads/2016/03/mem_map-6c.png)

struct page

include/linux/mm_types.h -1/4-

```

01  /*
02  * Each physical page in the system has a struct page associated with
03  * it to keep track of whatever it is we are using the page for at the
04  * moment. Note that we have no way to track which tasks are using
05  * a page, though if it is a pagecache page, rmap structures can tell us
06  * who is mapping it.
07  *
08  * If you allocate the page using alloc_pages(), you can use some of the
09  * space in struct page for your own purposes. The five words in the ma
10  in
11  * union are available, except for bit 0 of the first word which must be
12  * kept clear. Many users use this word to store a pointer to an object
13  * which is guaranteed to be aligned. If you use the same storage as
14  * page->mapping, you must restore it to NULL before freeing the page.
15  *
16  * If your page will not be mapped to userspace, you can also use the fo
17  ur
18  * bytes in the mapcount union, but you must call page_mapcount_reset()
19  * before freeing it.
20  *
21  * If you want to use the refcount field, it must be used in such a way
22  * that other CPUs temporarily incrementing and then decrementing the
23  * refcount does not cause problems. On receiving the page from
24  * alloc_pages(), the refcount will be positive.
25  *
26  * If you allocate pages of order > 0, you can use some of the fields
27  * in each subpage, but you may need to restore some of their values
28  * afterwards.
29  *
30  * SLUB uses cmpxchg_double() to atomically update its freelist and
31  * counters. That requires that freelist & counters be adjacent and
32  * double-word aligned. We align all struct pages to double-word
33  * boundaries, and ensure that 'freelist' is aligned within the

```



```

32 | * struct.
33 | */

01 | struct page {
02 |     unsigned long flags;                /* Atomic flags, some possibly
03 |                                         * updated asynchronously */
04 |     /*
05 |      * Five words (20/40 bytes) are available in this union.
06 |      * WARNING: bit 0 of the first word is used for PageTail(). That
07 |      * means the other users of this union MUST NOT use the bit to
08 |      * avoid collision and false-positive PageTail().
09 |      */
10 |     union {
11 |         struct {                        /* Page cache and anonymous pages */
12 |             /**
13 |              * @lru: Pageout list, eg. active_list protected
14 |              * lruvec->lru_lock. Sometimes used as a generi
15 |              * by the page owner.
16 |              */
17 |             struct list_head lru;
18 |             /* See page-flags.h for PAGE_MAPPING_FLAGS */
19 |             struct address_space *mapping;
20 |             pgoff_t index;              /* Our offset within map
21 |             ping. */
22 |             /**
23 |              * @private: Mapping-private opaque data.
24 |              * Usually used for buffer_heads if PagePrivate.
25 |              * Used for swp_entry_t if PageSwapCache.
26 |              * Indicates order in the buddy system if PageBu
27 |              */
28 |             unsigned long private;
29 |         };
30 |         struct {                        /* page_pool used by netstack */
31 |             /**
32 |              * @pp_magic: magic value to avoid recycling non
33 |              * page_pool allocated pages.
34 |              */
35 |             unsigned long pp_magic;
36 |             struct page_pool *pp;
37 |             unsigned long _pp_mapping_pad;
38 |             unsigned long dma_addr;
39 |             union {
40 |                 /**
41 |                  * dma_addr_upper: might require a 64-bi
42 |                  * value on 32-bit architectures.
43 |                  */
44 |                 unsigned long dma_addr_upper;
45 |                 /**
46 |                  * For frag page support, not supported
47 |                  * 32-bit architectures with 64-bit DMA.
48 |                  */
49 |                 atomic_long_t pp_frag_count;
50 |             };

```

First Word

- flags
 - Page Flags

1) Page cache or anonymous page

- `_lru`
 - Used to connect to an LRU list.
- `*mapping`
 - It contains pointers related to user mapping, and the bottom 2 bits are used as flags to distinguish their purpose.
 - When used as a page cache, it points to a `address_space` structure.
 - Non-LRU movable pages are managed by adding `PAGE_MAPPING_MOVABLE` flags to `address_space` structure pointers.
 - e.g. ZRAM, balloon driver
 - In the case of an anonymous page for users, it depends on whether the `CONFIG_KSM` kernel option is enabled.
 - When the KSM kernel option is not used, only the `PAGE_MAPPING_ANON(1)` flag is added and points to the anon mapping area, the `anon_vma` structure pointer.
 - If you are using the KSM kernel options, add the `PAGE_MAPPING_ANON(1)` and `PAGE_MAPPING_MOVABLE(2)` flags, and point to a pointer to a private struct for KSM.
- `index`
 - The offset value in the mapping area is contained.
- `private`
 - It contains private data used for mapping.
 - It is used for `buffer_heads` on private pages.
 - Swab is used for `swp_entry_t` of page cache.
 - The order of the buddy page is contained.

2) `page_pool` pages used by Netstack

- `pp_magic`
 - The purpose is to identify the use of network `page_pool`.
- `*pp`
 - `page_pool` a pointer to a struct.
- `_pp_mapping_pad`
 - padding is intended and is not actually used.
- `dma_addr`
 - DMA Address
- `dma_addr_upper`
 - When using a 32-bit DMA on a 64-bit system, the lower 32 bits of the DMA address are stored in the upper `dma_addr`, and the upper 32 bits are stored in this member.
- `pp_frag_count`
 - This is a frag counter managed by atomic.
 - It is not supported on 64-bit systems that use 32-bit DMA.

```

01      struct {          /* slab, slob and slub */
02          union {
03              struct list_head slab_list;      /* uses
lru */
04              struct {          /* Partial pages */
05                  struct page *next;
06 #ifdef CONFIG_64BIT
07                  int pages;      /* Nr of pages l
eft */
08                  int pobjects;   /* Approximate c
ount */
09 #else
10                  short int pages;
11                  short int pobjects;
12 #endif
13              };
14      };
15      kmem_cache *slab_cache; /* not slob */
16      /* Double-word boundary */
17      void *freelist; /* first free object */
18      union {
19          void *s_mem; /* slab: first object */
20          unsigned long counters; /* SLUB
*/
21          struct {          /* SLUB
*/
22              unsigned inuse:16;
23              unsigned objects:15;
24              unsigned frozen:1;
25          };
26      };
27  };
28  struct {          /* Tail pages of compound page */
29      unsigned long compound_head; /* Bit zero is s
et */
30
31      /* First tail page only */
32      unsigned char compound_dtor;
33      unsigned char compound_order;
34      atomic_t compound_mapcount;
35      unsigned int compound_nr; /* 1 << compound_order
*/
36  };
37  struct {          /* Second tail page of compound page */
38      unsigned long _compound_pad_1; /* compound_head
*/
39
40      atomic_t hpage_pinned_refcount;
41      /* For both global and memcg */
42      struct list_head deferred_list;

```

3) Slab, slob, slub pages

- slab_list
 - Used to connect to an LRU list.
- *next
 - Partial pages.
- pages
 - Contains a partial number of pages.
 - It contains the number of slub pages that are connected to the next, including itself.
- pobjects
 - Approximate number of objects

- It's not exact, but it contains the total number of free objects for the next slab page, including my slab page.
- This counter is often not calculated correctly due to the objects being free
- *slab_cache
 - Point to the associated slab cache.
- *freelist
 - This is a list of free objects waiting for you.
- *s_mem
 - slab's first object.
- counters
 - Use the following 32 bytes to access them at once.
 - inuse:16
 - The number of objects in use.
 - objects:15
 - Number of objects managed by slabs
 - frozen:1
 - It refers to whether it is a slab page that is managed by Per-CPU.

4) Compound tail pages

Note that the head page (page[0]) of the Compound page has a PG_head flag.

- compound_head
 - All tail pages that are not the header of the compound page contain the compound header page descriptor pointer, and set bit0 to 1.
- compound_dtors
 - On the first tail page (page[1]), put the compound page destructor identifier id.
 - Contains one of the following compound page destructor IDs:
 - NULL_COMPOUND_DTORS
 - COMPOUND_PAGE_DTORS
 - HUGETLB_PAGE_DTORS
 - TRANSHUGE_PAGE_DTORS
- compound_order
 - The first tail page (page[1]) contains the order of the compound page.
- compound_mapcount
 - The first tail page (page[1]) contains the number of mapping counts.
- _compound_pad_1
 - Do not use it.
- _compound_pad_2
 - Do not use it.
- deferred_list
 - It is used to detach a huge page from the second tail page (page[2]) and suspend it for a while instead of unmapping it immediately.

- Note: thp: introduce deferred_split_huge_page()
(<https://github.com/torvalds/linux/commit/9a982250f773cc8c76f1eee68a770b7cbf2faf78#diff-ff0276c422430aff67678952b177fe3d>)

include/linux/mm_types.h -3/4-

```

01      struct {          /* Page table pages */
02          unsigned long _pt_pad_1;          /* compound_head
03      */
04          pgtable_t pmd_huge_pte; /* protected by page->pt
05      l */
06          unsigned long _pt_pad_2;          /* mapping */
07          union {
08              struct mm_struct *pt_mm; /* x86 pgds onl
09              atomic_t pt_frag_refcount; /* powerpc */
10          };
11      #if ALLOC_SPLIT_PTLOCKS
12          spinlock_t *ptl;
13      #else
14          spinlock_t ptl;
15      #endif
16      };
17      struct {          /* ZONE_DEVICE pages */
18          /** @pgmap: Points to the hosting device page ma
19      p. */
20          struct dev_pagemap *pgmap;
21          void *zone_device_data;
22          /*
23      * ZONE_DEVICE private pages are counted as bein
24      * mapped so the next 3 words hold the mapping,
25      * and private fields from the source anonymous
26      * page cache page while the page is migrated to
27      * private memory.
28      * ZONE_DEVICE MEMORY_DEVICE_FS_DAX pages also
29      * use the mapping, index, and private fields wh
30      * pmem backed DAX files are mapped.
31      */
32      };
33      /** @rcu_head: You can use this to free a page by RCU.
34      */
35      struct rcu_head rcu_head;
36  };

```

5) Pages for page tables (PGD, PUD, PMD, PTE)

- _pt_pad_1
 - Do not use it.
- pmd_huge_pte
 - It contains a Page descriptor pointer that points to the PTE table for HUGE.
- _pt_pad_2
 - Do not use it.
- *pt_mm
 - x86 pgds only

- `*pt_frag_refcount`
 - powerpc only
- `*ptl` or `ptl`
 - The page table is a spinlock.
 - If the `spinlock_t` size is included in an unsigned long unit, use PTL, otherwise use `*PTL` to assign and point to `spinlock_t`.

6) Zone Devices Page

- `*pgmap`
 - Contains a `dev_pagemap` struct pointer for the zone device.
- `hmm_data`
 - It contains driver data for HMM device memory.
- `_zd_pad_1`
 - Do not use it.

include/linux/mm_types.h -4/4-

```

01         union {          /* This union is 4 bytes in size. */
02             /*
03             * If the page can be mapped to userspace, encodes the n
umber
04             * of times this page is referenced by a page table.
05             */
06             atomic_t _mapcount;
07
08             /*
09             * If the page is neither PageSlab nor mappable to users
pace,
10             * the value stored here may help determine what this pa
ge
11             * is used for. See page-flags.h for a list of page typ
es
12             * which are currently stored here.
13             */
14             unsigned int page_type;
15
16             unsigned int active;          /* SLAB */
17             int units;                   /* SLOB */
18         };
19
20         /* Usage count. *DO NOT USE DIRECTLY*. See page_ref.h */
21         atomic_t _refcount;
22
23 #ifdef CONFIG_MEMCG
24         unsigned long memcg_data;
25 #endif
26
27         /*
28         * On machines where all RAM is mapped into kernel address spac
e,
29         * we can simply calculate the virtual address. On machines with
30         * highmem some memory is mapped into kernel virtual memory
31         * dynamically, so we need a place to store that address.
32         * Note that this field could be 16 bits on x86 ... ☺
33         *
34         * Architectures with slow multiplication can define
35         * WANT_PAGE_VIRTUAL in asm/page.h

```

```

36      */
37  #if defined(WANT_PAGE_VIRTUAL)
38      void *virtual;                /* Kernel virtual address (NULL
    if                                not kmapped, ie. highmem) */
39
40  #endif /* WANT_PAGE_VIRTUAL */
41
42  #ifdef LAST_CPUPID_NOT_IN_PAGE_FLAGS
43      int _last_cpupid;
44  #endif
45  } _struct_page_alignment;

```

Other-1

The `_mapcount`, `page_type`, `active`, and `units` below are defined as union types that are selected and used according to the page type.

- `_mapcount`
 - Mapping counts used by users
 - Each time it is mapped to multiple user spaces, the counter is incremented.
 - If a shared page is shared by multiple user processes, it will be multi-mapped, and the mapping counter will be incremented by the number.
- `page_type`
 - Except for the user mapping page and the slab page, the rest of the kernel pages specify the page type and store and use the following flags.
 - `PG_buddy`
 - `PG_ballon`
 - `PG_kmemcg`
 - `PG_table`
- `active`
 - slab: Used in the allocator.
- `units`
 - slob: Used in the allocator.

Other-2

- `_refcount`
 - The reference counter should not be accessed directly by the user.
- `memcg_data`
 - When using a memory cgroup on a system with `CONFIG_MEMCG` kernel options, it contains a `obj_cgroup` struct pointer that can be used, and depending on the type of allocation, it uses the following two flags:
 - `MEMCG_DATA_OBJCGS`
 - `MEMCG_DATA_KMEM`
- `*virtual`
 - If you use the `WANT_PAGE_VIRTUAL` kernel option, in some 32-bit architectures that use highmem, if the highmem page is mapped to the kernel address space, it will contain the mapped kernel virtual address.
 - In the case of ARM32, `HASHED_PAGE_VIRTUAL` is used.
 - When mapping a highmem page, we don't use this member because we use a separate allocation of `page_address_map` structs.

- 64-bit systems use highmem, so you don't need to use this member.
- `_last_cpupid`
 - There are times when you need to include the CPU and PID in the flags member of the page, but on a 32-bit system you can't include both in 32-bit. In this case, use the `LAST_CPUPID_NOT_IN_PAGE_FLAGS` kernel option to store the last used CPU and PID in this member. They are used for NUMA balancing.
 - 64-bit systems use the flags member to store the last CPU and PID information, so there is no need to use this member.

consultation

- Memory Model -1- (Basic) (<http://jake.dothome.co.kr/mm-1>) | 문c
 - Memory Model -2- (mem_map) (http://jake.dothome.co.kr/mem_map) | Sentence C – Current post
 - Memory Model -3- (Sparse Memory) (<http://jake.dothome.co.kr/sparsemem/>) | 문c
 - Memory Model -4- (APIs) (http://jake.dothome.co.kr/mem_map_page) | 문c
 - ZONE Type (<http://jake.dothome.co.kr/zone-types>) | Qc
 - `bootmem_init` (http://jake.dothome.co.kr/bootmem_init-64) | Qc
 - `zone_sizes_init()` (http://jake.dothome.co.kr/free_area_init_node/) | Qc
 - NUMA -1- (ARM64 initialization) (<http://jake.dothome.co.kr/numa-1>) | Qc
 - `build_all_zonelists()` (http://jake.dothome.co.kr/build_all_zonelists) | Qc
-
- Memory: the flat, the discontiguous, and the sparse(2019) (<https://lwn.net/Articles/789304/>) | LWN.net
 - Generic page-pool recycle facility? (2016) | Jesper Dangaard Brouer – Download pdf (http://people.netfilter.org/hawk/presentations/MM-summit2016/generic_page_pool_mm_summit2016.pdf)

6 thoughts to “Memory Model -2- (mem_map)”



2019-04-12 18:10 (http://jake.dothome.co.kr/mem_map/#comment-208486)

I don't think I understood it correctly.

When you know the page structure, how do you know the physical memory that the page structure points to, i.e. the page frame#는

RESPONSE (/MEM_MAP/?REPLYTOCOM=208486#RESPOND)



MOON YOUNG-IL (HTTP://JAKE.DOTHOME.CO.KR)

2019-04-15 11:25 (http://jake.dothome.co.kr/mem_map/#comment-208628)

Hello? Moon Young-il of Munc Blog.

You can use the `page_to_pfn()` API to get the page frame number (pfn) from a pointer to the page structure pointing to that page.

An API that does the opposite is the `pfn_to_page()` API.

Keep up the good work.

RESPONSE (/MEM_MAP/?REPLYTOCOM=208628#RESPOND)



2020-05-17 09:00 (http://jake.dothome.co.kr/mem_map/#comment-249307)

Hello good article, I am reading it carefully.

There is no MM-7B, MM-9A picture.

RESPONSE (/MEM_MAP/?REPLYTOCOM=249307#RESPOND)



MOON YOUNG-IL (HTTP://JAKE.DOTHOME.CO.KR)

2020-06-22 11:18 (http://jake.dothome.co.kr/mem_map/#comment-257080)

I've included the picture again.

Thanks for letting us know.

RESPONSE (/MEM_MAP/?REPLYTOCOM=257080#RESPOND)



2020-09-22 20:08 (http://jake.dothome.co.kr/mem_map/#comment-277674)

Hello. Do you have any idea how memory is initialized when Linux boots? (mem_map arrays, etc.) Even if I tried to allocate an array like that, I'd need a memory allocation policy. I'm curious how to initialize it at the very beginning.

RESPONSE (/MEM_MAP/?REPLYTOCOM=277674#RESPOND)



MOON YOUNG-IL (HTTP://JAKE.DOTHOME.CO.KR)

2020-09-24 07:00 (http://jake.dothome.co.kr/mem_map/#comment-278068)

안녕하세요?

2024/1/1 14:02

Memory Model -2-(mem_map) - Munc Blog

리눅스 부팅 후 사용되는 대표적인 메모리 할당 시스템은 buddy, slab, per-cpu, cma, ... 등이 있습니다.

이를 초기화하여 운용하기 전에는 memblock을 사용하여 영역 정보만을 관리하고 있습니다.

메모리 할당자를 사용하기 전에 memblock을 통해 미리 필요한 영역을 reserve하여 등록하고, 사용할 때에는 원시적으로 직접 해당 주소 영역에 직접 기록하는 방법을 사용합니다. 그런 후 추후 페이지 할당자인 buddy 시스템을 준비할 때 그 reserve 영역을 피한 나머지 메모리 영역만 buddy 시스템에 free 페이지로 등록하여 사용합니다.

감사합니다.

응답 (/MEM_MAP/?REPLYTOCOM=278068#RESPOND)

댓글 남기기

이메일은 공개되지 않습니다. 필수 입력창은 * 로 표시되어 있습니다

댓글

이름 *

이메일 *

웹사이트

댓글 작성

◀ Fixmap (<http://jake.dothome.co.kr/fixmap/>)

페이지 테이블 API - ARM32 ▶ (<http://jake.dothome.co.kr/pt-api/>)