

# NUMA -3- (Memory policy)

📅 2019-07-25 (<http://jake.dothome.co.kr/numa-3-memory-policy/>) 👤 Moon Young-il

(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

<kernel v5.0>

## NUMA -3- (Memory policy)

### Memory Policy

There are the following types of NUMA memory policies:

- MPOL\_DEFAULT
  - This mode is only used inside the Memory Policy API. NULL and fallback internally to use the system's default memory policy.
- MPOL\_PREFERRED
  - Specify and assign one preferred node. If you run out of memory, you can allocate it on another node.
  - However, if it is used with the MPOL\_F\_LOCAL flag, the preferred function is ignored and the local node allocation is given priority.
- MPOL\_BIND
  - Allocate memory only on the specified bind nodes. If you run out of memory, you can't allocate it on other nodes.
- MPOL\_INTERLEAVE
  - Round Robin at the specified interleaved nodes. If you run out of memory, you can allocate it on another node.
- MPOL\_LOCAL
  - Change the mode to MPOL\_PREFERRED, and use it with the MPOL\_F\_LOCAL flag to give preference to local nodes.

### Memory Policy Flags

Use the following various memory policy-related flags:

- Flags used in set\_mempolicy()
  - MPOL\_F\_STATIC\_NODES(0x8000)
    - Specifying a static node
  - MPOL\_F\_RELATIVE\_NODES(0x4000)
    - Specifying Relative Nodes
- Flags used in get\_mempolicy()

- MPOL\_F\_NODE(1)
  - Returning the next IL node instead of the node masque
- MPOL\_F\_ADDR(2)
  - Search for VMA by address
- MPOL\_F\_MEMS\_ALLOWED(4)
  - return allowed memories
- Flags used in mbind()
  - MPOL\_MF\_STRICT(1)
  - MPOL\_MF\_MOVE(2)
  - MPOL\_MF\_MOVE\_ALL(4)
  - MPOL\_MF\_LAZY(8)
  - MPOL\_MF\_INTERNAL(16)
- Internal flags used with mods
  - MPOL\_F\_SHARED(1)
    - Sharing Policies
  - MPOL\_F\_LOCAL(2)
    - Preferred Local Node Assignment
  - MPOL\_F\_MOF(8)
    - Migrate on fault
  - MPOL\_F\_MORON(16)
    - Migrate On protnone Reference On Node

The following shows that 20 CPU cores are used on each of the two nodes.

```
01 $ numactl --hardware
02 available: 2 nodes (0-1)
03 node 0 cpus: 0 1 2 3 4 5 6 7 8 9 20 21 22 23 24 25 26 27 28 29
04 node 0 size: 32654 MB
05 node 0 free: 18259 MB
06 node 1 cpus: 10 11 12 13 14 15 16 17 18 19 30 31 32 33 34 35 36 37 38 39
07 node 1 size: 32768 MB
08 node 1 free: 15491 MB
09 node distances:
10 node    0    1
11   0:   10   21
12   1:   21   10
```

The following shows how the Numa node policy uses default.

```
1 $ numactl --show
2 policy: default
3 preferred node: current
4 physcpubind: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22
  23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39
5 cpubind: 0 1
6 nodebind: 0 1
7 membind: 0 1
```

# NUMA Memory Policy

## default\_policy Global Objects

mm/mempolicy.c

```

1  /*
2   * run-time system-wide default policy => local allocation
3   */
4  static struct mempolicy default_policy = {
5      .refcnt = ATOMIC_INIT(1), /* never free it */
6      .mode = MPOL_PREFERRED,
7      .flags = MPOL_F_LOCAL,
8  };

```

Specify the memory policy to prioritize local nodes.

## Knowing the memory policy for a task

### get\_task\_policy()

mm/mempolicy.c

```

01 struct mempolicy *get_task_policy(struct task_struct *p)
02 {
03     struct mempolicy *pol = p->mempolicy;
04     int node;
05
06     if (pol)
07         return pol;
08
09     node = numa_node_id();
10     if (node != NUMA_NO_NODE) {
11         pol = &preferred_node_policy[node];
12         /* preferred_node_policy is not initialised early in boo
13         t */
14         if (pol->mode)
15             return pol;
16     }
17     return &default_policy;
18 }

```

Returns the memory policy of the current task according to the following case:

- If the task has a policy specified, then -> 1) The memory policy specified in the task
- If the specified node exists, > 2) Use the specified node-first memory policy
- -> on request without a specified node 3) System Default Memory Policy (Local Node Preferred)

### policy\_node()

mm/mempolicy.c

```

01 /* Return the node id preferred by the given mempolicy, or the given id
02  */
03 static int policy_node(gfp_t gfp, struct mempolicy *policy, int nd)
04 {
05     if (policy->mode == MPOL_PREFERRED && !(policy->flags & MPOL_F_L
06     OCAL))
07         nd = policy->v.preferred_node;

```

```

07     else {
08         /*
09         * __GFP_THISNODE shouldn't even be used with the bind p
10         * because we might easily break the expectation to stay
11         * on the requested node and not break the policy.
12         */
13         WARN_ON_ONCE(policy->mode == MPOL_BIND && (gfp & __GFP_T
14         HISNODE));
15     }
16     return nd;
17 }

```

If the memory policy is in Preferred mode and a Preferred node is specified, it returns the node number. If not, it returns an input argument @nd.

- If preferred requests to use a local node, it will simply return @nd.

## policy\_nodemask()

mm/mempolicy.c

```

1  /*
2  * Return a nodemask representing a mempolicy for filtering nodes for
3  * page allocation
4  */

01 static nodemask_t *policy_nodemask(gfp_t gfp, struct mempolicy *policy)
02 {
03     /* Lower zones don't get a nodemask applied for MPOL_BIND */
04     if (unlikely(policy->mode == MPOL_BIND) &&
05         apply_policy_zone(policy, gfp_zone(gfp)) &&
06         cpuset_nodemask_valid_mems_allowed(&policy->v.no
07     des))
08         return &policy->v.nodes;
09     return NULL;
10 }

```

## Interleaved Memory Policy

### interleave\_nodes()

mm/mempolicy.c

```

01 /* Do dynamic interleaving for a process */
02 static unsigned interleave_nodes(struct mempolicy *policy)
03 {
04     unsigned nid, next;
05     struct task_struct *me = current;
06
07     nid = me->il_next;
08     next = next_node(nid, policy->v.nodes);
09     if (next >= MAX_NUMNODES)
10         next = first_node(policy->v.nodes);
11     if (next < MAX_NUMNODES)
12         me->il_next = next;
13     return nid;
14 }

```

Traversal the interleaved nodes and returns the node number.

- currnet->il\_next In the meantime, keep track of the node you want to assign next in current->il\_next.
- In il\_next, nodes to be used when the memory policy is set to MPOL\_INTERLEAVE are assigned by the interleave (round robin) method.
  - next\_node()
    - Finds the next node of the specified node for the node bitmap. If it can't find it, it returns a MAX\_NUMNODES.
  - first\_node()
    - Node: Knows the node located at the beginning of the bitmap.

## alloc\_page\_interleave()

mm/mempolicy.c

```

1  | /* Allocate a page in interleaved policy.
2  |    Own path because it needs to do special accounting. */

01 | static struct page *alloc_page_interleave(gfp_t gfp, unsigned order,
02 |                                           unsigned nid)
03 | {
04 |     struct page *page;
05 |
06 |     page = __alloc_pages(gfp, order, nid);
07 |     /* skip NUMA_INTERLEAVE_HIT counter update if numa stats is disabled */
08 |     if (!static_branch_likely(&vm_numa_stat_key))
09 |         return page;
10 |     if (page && page_to_nid(page) == nid) {
11 |         preempt_disable();
12 |         __inc_numa_state(page_zone(page), NUMA_INTERLEAVE_HIT);
13 |         preempt_enable();
14 |     }
15 |     return page;
16 | }
```

The interleave memory policy allocates  $2^{\text{order}}$  pages of contiguous physical memory.

- The gfp flag requested in line 6 of the code, the node allocates  $2^{\text{order}}$  pages of contiguous physical memory.
  - Depending on the gfp flag, you can either get a full node zonelist of node\_zonelist[0], or a node\_zonelist[1] containing only the specified node zones.
- If you are not using the NUMA statistic in line 8~9 of the code, skip it.
- Update the NUMA statistics on lines 10~14 of the code.
  - Only if the page is assigned from the requested node will the NUMA\_INTERLEAVE\_HIT stat of the zone be incremented.

## Main Structure

### mempolicy struct

include/linux/mempolicy.h

```

01  /*
02  * Describe a memory policy.
03  *
04  * A mempolicy can be either associated with a process or with a VMA.
05  * For VMA related allocations the VMA policy is preferred, otherwise
06  * the process policy is used. Interrupts ignore the memory policy
07  * of the current process.
08  *
09  * Locking policy for interlave:
10  * In process context there is no locking because only the process acces
11  * its own state. All vma manipulation is somewhat protected by a down_r
12  * mmap_sem.
13  *
14  * Freeing policy:
15  * Mempolicy objects are reference counted. A mempolicy will be freed w
16  * mpol_put() decrements the reference count to zero.
17  *
18  * Duplicating policy objects:
19  * mpol_dup() allocates a new mempolicy and copies the specified mempoli
20  * to the new storage. The reference count of the new object is initial
21  * to 1, representing the caller of mpol_dup().
22  */

01  struct mempolicy {
02      atomic_t refcnt;
03      unsigned short mode; /* See MPOL_* above */
04      unsigned short flags; /* See set_mempolicy() MPOL_F_* above */
05      union {
06          short preferred_node; /* preferred */
07          nodemask_t nodes; /* interleave/bind */
08          /* undefined for default */
09      } v;
10      union {
11          nodemask_t cpuset_mems_allowed; /* relative to these nod
12          nodemask_t user_nodemask; /* nodemask passed by us
13      } w;
14  };

```



## consultation

- NUMA with Linux (<https://lunatime.net/2016/07/14/numa-with-linux/>) | Lunatine's Box
- Local and Remote Memory: Memory in a Linux/NUMA System | Christoph Lameter – pdf 다운로드 ([http://gentwo.org/christoph/Presentations-2000-2009/2006-OLS-Local\\_and\\_Remote\\_Memory.pdf](http://gentwo.org/christoph/Presentations-2000-2009/2006-OLS-Local_and_Remote_Memory.pdf))
- NUMA Best Practices for Dell PowerEdge 12th Generation Servers | Dell – pdf 다운로드 ([https://lafibre.info/images/materiel/201212\\_dell\\_bios\\_tuning\\_for\\_performance.pdf](https://lafibre.info/images/materiel/201212_dell_bios_tuning_for_performance.pdf))
- What is Linux Memory Policy? ([https://www.kernel.org/doc/Documentation/vm/numa\\_memory\\_policy.txt](https://www.kernel.org/doc/Documentation/vm/numa_memory_policy.txt)) | kernel.org
- NUMA API for Linux (<https://lwn.net/Articles/79100/>) | LWN.net
- numa – overview of Non-Uniform Memory Architecture (<http://man7.org/linux/man-pages/man7/numa.7.html>) | man7.org

LEAVE A COMMENT

Your email will not be published. Required fields are marked with \*

Comments

name \*

email \*

Website

WRITE A COMMENT

◀ Debug Memory-4- (Idle Page Trace) (<http://jake.dothome.co.kr/debug-mem-4/>)

kernel/head. S – ARM64 (old for v5.0) ▶ (<http://jake.dothome.co.kr/head-64/>)