

# Per-cpu -4- (atomic operations)

📅 2016-06-01 (<http://jake.dothome.co.kr/per-cpu-atomic/>) 👤 Moon Young-il

(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

## Per-cpu -4- (atomic operations)

### this\_cpu\_cmpxchg\_double()

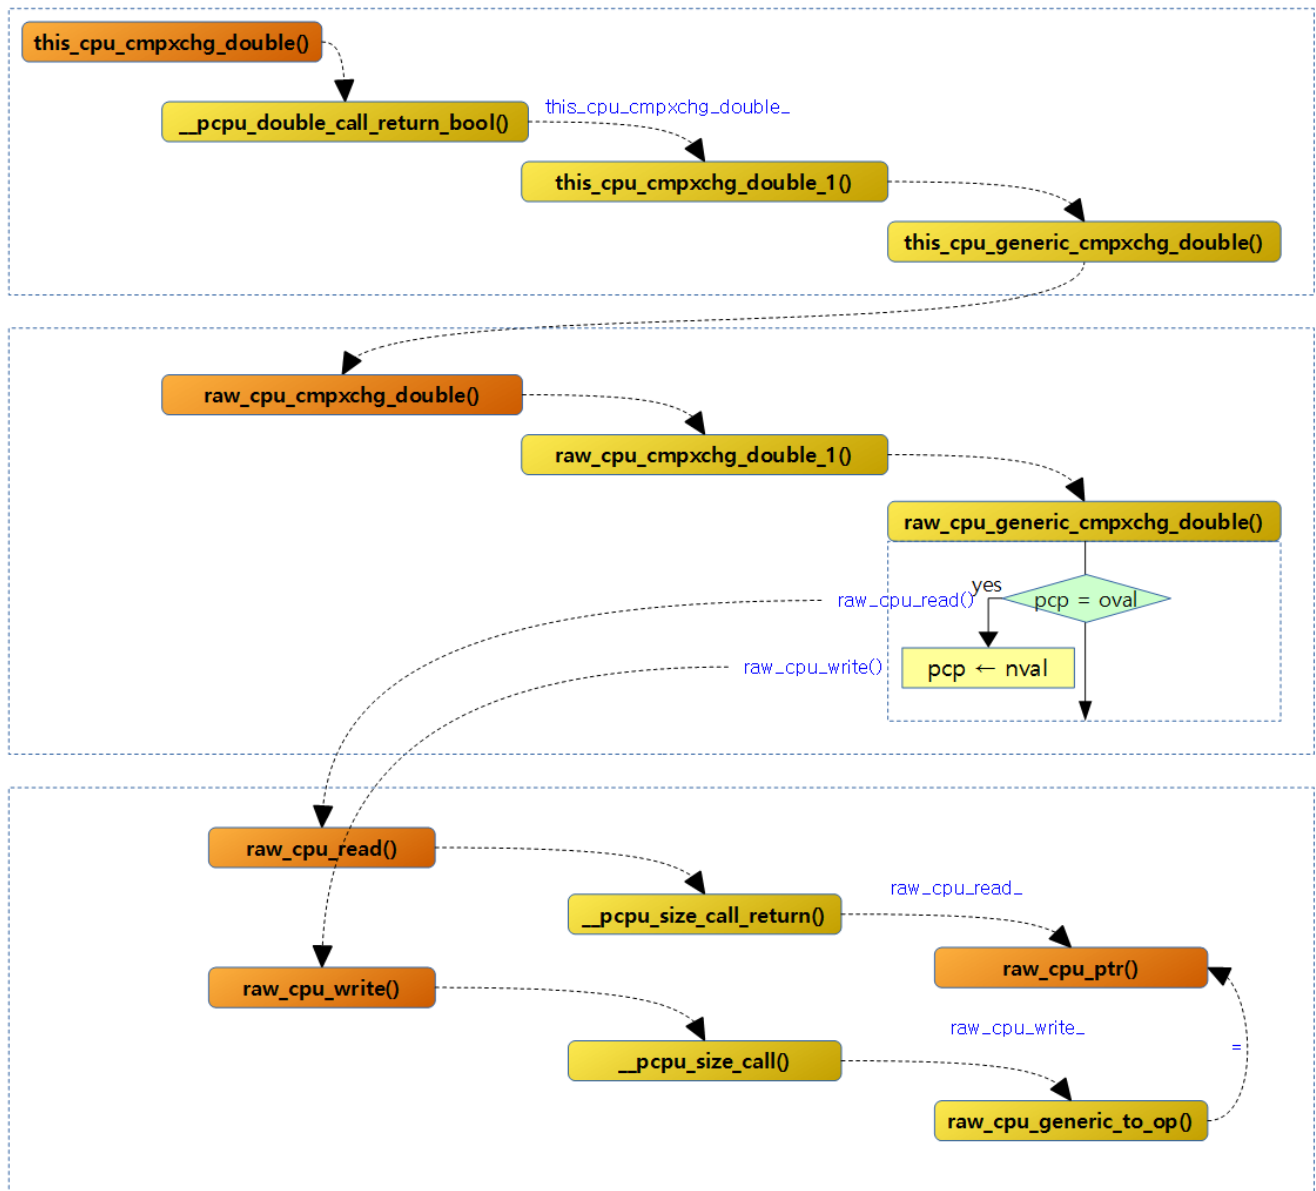
include/linux/percpu-defs.h

```
1 | #define this_cpu_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, nval2)  
  | \  
2 |     __pcpu_double_call_return_bool(this_cpu_cmpxchg_double_, pcp1, p  
  | cp2, oval1, oval2, nval1, nval2)
```

If the per-CPU value of PCP(PCP1, PCP2) is the same as the old value (oval1, oval2), then the new value (nval1, nval2) is atomically substituted into the pcp.

- How it differs from the cmpxchg\_double() function
  - Expect faster atomic operations to replace per-CPU values that don't need to compete with other CPUs.
    - In the ARM architecture, local irq is blocked only during atomic operation.
    - In the arm64 architecture, preemption is prevented only during atomic operations.

The figure below shows how the this\_cpu\_cmpxchg\_double() function is processed.



## \_\_pcpu\_double\_call\_return\_bool()

include/linux/percpu-defs.h

```

01  /*
02  * Special handling for cmpxchg_double.  cmpxchg_double is passed two
03  * percpu variables.  The first has to be aligned to a double word
04  * boundary and the second has to follow directly thereafter.
05  * We enforce this on all architectures even if they don't support
06  * a double cmpxchg instruction, since it's a cheap requirement, and it
07  * avoids breaking the requirement for architectures with the instructio
n.
08  */
09  #define __pcpu_double_call_return_bool(stem, pcp1, pcp2, ...)
10  ({
11      \
12      bool pdcrb_ret__;
13      \
14      __verify_pcpu_ptr(&(pcp1));
15      \
16      BUILD_BUG_ON(sizeof(pcp1) != sizeof(pcp2));
17      \
18      VM_BUG_ON((unsigned long)(amp;(pcp1)) % (2 * sizeof(pcp1)));
19      \
20  })

```

```

15 VM_BUG_ON((unsigned long)(ampcp2)) !=
16     (unsigned long)(ampcp1) + sizeof(pcp1));
17 switch(sizeof(pcp1)) {
18 case 1: pdcrb_ret__ = stem##1(pcp1, pcp2, __VA_ARGS__); break;
19 case 2: pdcrb_ret__ = stem##2(pcp1, pcp2, __VA_ARGS__); break;
20 case 4: pdcrb_ret__ = stem##4(pcp1, pcp2, __VA_ARGS__); break;
21 case 8: pdcrb_ret__ = stem##8(pcp1, pcp2, __VA_ARGS__); break;
22 default:
23     __bad_size_call_parameter(); break;
24 }
25 pdcrb_ret__;
26 })

```

If the value of PCP(PCP1, PCP2) is the same as the old value (oval1, oval2), substitute the new value (nval1, nval2) for the pcp.

- Depending on the data length, the argument will call stem1~stem8.
  - e.g. stem=this\_cpu\_cmpxchg\_double\_
    - this\_cpu\_cmpxchg\_double\_1, this\_cpu\_cmpxchg\_double\_2, this\_cpu\_cmpxchg\_double\_4, this\_cpu\_cmpxchg\_double\_8

### **this\_cpu\_cmpxchg\_double\_1()**

include/asm-generic/percpu.h

```

01 #ifndef this_cpu_cmpxchg_double_1
02 #define this_cpu_cmpxchg_double_1(pcp1, pcp2, oval1, oval2, nval1, nval2) \
03     this_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, \
04     nval2)
05 #endif
06 #ifndef this_cpu_cmpxchg_double_2
07 #define this_cpu_cmpxchg_double_2(pcp1, pcp2, oval1, oval2, nval1, nval2) \
08     this_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, \
09     nval2)
10 #endif
11 #ifndef this_cpu_cmpxchg_double_4
12 #define this_cpu_cmpxchg_double_4(pcp1, pcp2, oval1, oval2, nval1, nval2) \
13     this_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, \
14     nval2)
15 #endif
16 #ifndef this_cpu_cmpxchg_double_8
17 #define this_cpu_cmpxchg_double_8(pcp1, pcp2, oval1, oval2, nval1, nval2) \
18     this_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, \
19     nval2)
20 #endif

```

If the value of PCP(PCP1, PCP2) is the same as the old value (oval1, oval2), substitute the new value (nval1, nval2) for the pcp.

- On a 32-bit arm, there is no operation to atomically process a double word, so it calls the generic function.

### **this\_cpu\_generic\_cmpxchg\_double()**

include/asm-generic/percpu.h

```

01 | #define this_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
02 |   ({
03 |     int __ret;
04 |     unsigned long __flags;
05 |     raw_local_irq_save(__flags);
06 |     __ret = raw_cpu_generic_cmpxchg_double(pcp1, pcp2,
07 |                                           oval1, oval2, nval1, nval2);
08 |     raw_local_irq_restore(__flags);
09 |     __ret;
10 | })

```

If the value of PCP(PCP1, PCP2) is the same as the old value (oval1, oval2) while the interrupt is disabled for a while, the new value (NVL1, NVL2) is substituted for the PCP, and the interrupt is reversed.

### **raw\_cpu\_cmpxchg\_double()**

include/linux/percpu-defs.h

```

1 | #define raw_cpu_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1, nval2) \
2 |   __pcpu_double_call_return_bool(raw_cpu_cmpxchg_double_, pcp1, pc
   |   p2, oval1, oval2, nval1, nval2)

```

If the value of pcp(pcp1, pcp2) is equal to the old value (oval1, oval2), it returns true if it succeeds in substituting a new value (nval1, nval2) into the pcp.

### **raw\_cpu\_cmpxchg\_double\_1()**

include/asm-generic/percpu.h

```

01 | #ifndef raw_cpu_cmpxchg_double_1
02 | #define raw_cpu_cmpxchg_double_1(pcp1, pcp2, oval1, oval2, nval1, nval2) \
03 |     raw_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
   |     nval2)
04 | #endif
05 | #ifndef raw_cpu_cmpxchg_double_2
06 | #define raw_cpu_cmpxchg_double_2(pcp1, pcp2, oval1, oval2, nval1, nval2) \
   |

```

```

07     raw_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
08     nval2)
08 #endif
09 #ifndef raw_cpu_cmpxchg_double_4
10 #define raw_cpu_cmpxchg_double_4(pcp1, pcp2, oval1, oval2, nval1, nval2)
11     \
12     raw_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
13     nval2)
12 #endif
13 #ifndef raw_cpu_cmpxchg_double_8
14 #define raw_cpu_cmpxchg_double_8(pcp1, pcp2, oval1, oval2, nval1, nval2)
15     \
16     raw_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
17     nval2)
16 #endif

```

If the value of pcp(pcp1, pcp2) is equal to the old value (oval1, oval2), it returns true if it succeeds in substituting a new value (nval1, nval2) into the pcp.

- On a 32-bit arm, there is no operation to atomically process a double word, so it calls the generic function.

### **raw\_cpu\_generic\_cmpxchg\_double()**

include/asm-generic/percpu.h

```

01 #define raw_cpu_generic_cmpxchg_double(pcp1, pcp2, oval1, oval2, nval1,
02 nval2) \
03 ({
04     \
05     int __ret = 0;
06     \
07     if (raw_cpu_read(pcp1) == (oval1) &&
08         raw_cpu_read(pcp2) == (oval2)) {
09         \
10         raw_cpu_write(pcp1, nval1);
11         raw_cpu_write(pcp2, nval2);
12         \
13         __ret = 1;
14     }
15     \
16     (__ret);
17 })

```

If the value of pcp(pcp1, pcp2) is equal to the old value (oval1, oval2), it returns true if it succeeds in substituting a new value (nval1, nval2) into the pcp.

### **raw\_cpu\_read()**

include/linux/percpu-defs.h

```

1 #define raw_cpu_read(pcp)                __pcpu_size_call_return(raw_cpu_
    read_, pcp)

```

pcp value.

## raw\_cpu\_write()

include/linux/percpu-defs.h

```
1 | #define raw_cpu_write(pcp, val)          __pcpu_size_call(raw_cpu_write_,
    | pcp, val)
```

Substitute the val value for the pcp value.

## \_\_pcpu\_size\_call\_return()

include/linux/percpu-defs.h

```
01 | #define __pcpu_size_call_return(stem, variable)
    | \
02 | ({
    | \
03 |     typeof(variable) pscr_ret__;
    | \
04 |     __verify_pcpu_ptr(&(variable));
    | \
05 |     switch(sizeof(variable)) {
    | \
06 |     case 1: pscr_ret__ = stem##1(variable); break;
    | \
07 |     case 2: pscr_ret__ = stem##2(variable); break;
    | \
08 |     case 4: pscr_ret__ = stem##4(variable); break;
    | \
09 |     case 8: pscr_ret__ = stem##8(variable); break;
    | \
10 |     default:
    | \
11 |         __bad_size_call_parameter(); break;
    | \
12 |     }
    | \
13 |     pscr_ret__;
    | \
14 | })
```

According to the size of the variable type, call stem1~8 specified as an argument and return the pcp value.

- e.g. stem=raw\_cpu\_read\_
  - voice=raw\_cpu\_read\_1, vote=raw\_cpu\_read\_2, voice=raw\_cpu\_read\_4,
   
vote=raw\_cpu\_read\_8

## \_\_pcpu\_size\_call()

include/linux/percpu-defs.h

```
01 | #define __pcpu_size_call(stem, variable, ...)
    | \
02 | do {
    | \
```

```

03 |     __verify_pcpu_ptr(&(variable));
04 |     \
05 |     switch(sizeof(variable)) {
06 |         \
07 |         case 1: stem##1(variable, __VA_ARGS__);break;
08 |         \
09 |         case 2: stem##2(variable, __VA_ARGS__);break;
10 |         \
11 |         case 4: stem##4(variable, __VA_ARGS__);break;
12 |         \
13 |         case 8: stem##8(variable, __VA_ARGS__);break;
14 |         \
15 |         default:
16 |             \
17 |             __bad_size_call_parameter();break;
18 |         \
19 |     }
20 | } while (0)

```

Call STEM1~8 specified as an argument according to the size of the variable type and assign the value to the PCP value.

- e.g. stem=raw\_cpu\_write\_
  - voice=raw\_cpu\_write\_1, vote=raw\_cpu\_write\_2, voice=raw\_cpu\_write\_4, vote=raw\_cpu\_write\_8

### raw\_cpu\_read\_1()

include/asm-generic/percpu.h

```

01 | #ifndef raw_cpu_read_1
02 | #define raw_cpu_read_1(pcp) (*raw_cpu_ptr(&(pcp)))
03 | #endif
04 | #ifndef raw_cpu_read_2
05 | #define raw_cpu_read_2(pcp) (*raw_cpu_ptr(&(pcp)))
06 | #endif
07 | #ifndef raw_cpu_read_4
08 | #define raw_cpu_read_4(pcp) (*raw_cpu_ptr(&(pcp)))
09 | #endif
10 | #ifndef raw_cpu_read_8
11 | #define raw_cpu_read_8(pcp) (*raw_cpu_ptr(&(pcp)))
12 | #endif

```

Read the value of pcp.

### raw\_cpu\_write\_1()

include/asm-generic/percpu.h

```

01 | #ifndef raw_cpu_write_1
02 | #define raw_cpu_write_1(pcp, val) raw_cpu_generic_to_op(pcp, val,
03 | =)
04 | #endif
05 | #ifndef raw_cpu_write_2
06 | #define raw_cpu_write_2(pcp, val) raw_cpu_generic_to_op(pcp, val,
07 | =)
08 | #endif
09 | #ifndef raw_cpu_write_4
10 | #define raw_cpu_write_4(pcp, val) raw_cpu_generic_to_op(pcp, val,
11 | =)
12 | #endif

```

```

09 | #endif
10 | #ifndef raw_cpu_write_8
11 | #define raw_cpu_write_8(pcp, val)      raw_cpu_generic_to_op(pcp, val,
    | =)
12 | #endif

```

Store the val value in pcp.

### **raw\_cpu\_generic\_to\_op()**

include/asm-generic/percpu.h

```

1 | #define raw_cpu_generic_to_op(pcp, val, op)
  | \
2 | do {
  | \
3 |     *raw_cpu_ptr(&(pcp)) op val;
  | \
4 | } while (0)

```

OP operations on pcp and val values.

- e.g. raw\_cpu\_generic\_to\_op(pcp, val, +=)
  - pcp += val

## **consultation**

- per-cpu -1- (Basic) (<http://jake.dothome.co.kr/per-cpu>) | 문c
- per-cpu -2- (initialize) ([http://jake.dothome.co.kr/setup\\_per\\_cpu\\_areas](http://jake.dothome.co.kr/setup_per_cpu_areas)) | Qc
- per-cpu -3- (dynamic allocation (<http://jake.dothome.co.kr/per-cpu-dynamic>)) | Qc
- Per-cpu -4- (atomic operations) | Question C – Current Article

---

### **LEAVE A COMMENT**

Your email will not be published. Required fields are marked with \*

Comments

name \*



email \*

Website

WRITE A COMMENT

◀ Zonned Allocator -3- (Buddy Page Allocation) (<http://jake.dothome.co.kr/buddy-alloc/>)

Slub Memory Allocator -13- (slabinfo) ▶ (<http://jake.dothome.co.kr/slub-slabinfo/>)

Munc Blog (2015 ~ 2024)