

Slub Memory Allocator -9- (cache shrink)

📅 2016-06-07 (<http://jake.dothome.co.kr/slub-cache-shrink/>) 👤 Moon Young-il

(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

<kernel v5.0>

kmem_cache_shrink()

mm/slab_common.c

```

1  /**
2   * kmem_cache_shrink - Shrink a cache.
3   * @cachep: The cache to shrink.
4   *
5   * Releases as many slabs as possible for a cache.
6   * To help debugging, a zero exit status indicates all slabs were releas
7   ed.
8   */
9
10 int kmem_cache_shrink(struct kmem_cache *cachep)
11 {
12     int ret;
13
14     get_online_cpus();
15     get_online_mems();
16     ret = __kmem_cache_shrink(cachep);
17     put_online_mems();
18     put_online_cpus();
19     return ret;
20 }
21 EXPORT_SYMBOL(kmem_cache_shrink);

```

Find the revocable slab pages in the requested slab cache and cancel all of them.

__kmem_cache_shrink()

mm/slub.c

```

1  /**
2   * kmem_cache_shrink discards empty slabs and promotes the slabs filled
3   * up most to the head of the partial lists. New allocations will then
4   * fill those up and thus they can be removed from the partial lists.
5   *
6   * The slabs with the least items are placed last. This results in them
7   * being allocated from last increasing the chance that the last objects
8   * are freed in them.
9   */
10
11 int __kmem_cache_shrink(struct kmem_cache *s)
12 {
13     int node;
14     int i;
15     struct kmem_cache_node *n;
16     struct page *page;
17     struct page *t;
18     struct list_head discard;

```

```

09 struct list_head promote[SHRINK_PROMOTE_MAX];
10 unsigned long flags;
11 int ret = 0;
12
13 flush_all(s);
14 for_each_kmem_cache_node(s, node, n) {
15     INIT_LIST_HEAD(&discard);
16     for (i = 0; i < SHRINK_PROMOTE_MAX; i++)
17         INIT_LIST_HEAD(promote + i);
18
19     spin_lock_irqsave(&n->list_lock, flags);
20
21     /*
22      * Build lists of slabs to discard or promote.
23      *
24      * Note that concurrent frees may occur while we hold th
25      * list_lock. page->inuse here is the upper limit.
26      */
27     list_for_each_entry_safe(page, t, &n->partial, lru) {
28         int free = page->objects - page->inuse;
29
30         /* Do not reread page->inuse */
31         barrier();
32
33         /* We do not keep full slabs on the list */
34         BUG_ON(free <= 0);
35
36         if (free == page->objects) {
37             list_move(&page->lru, &discard);
38             n->nr_partial--;
39         } else if (free <= SHRINK_PROMOTE_MAX)
40             list_move(&page->lru, promote + free -
41 1);
42     }
43
44     /*
45      * Promote the slabs filled up most to the head of the
46      * partial list.
47      */
48     for (i = SHRINK_PROMOTE_MAX - 1; i >= 0; i--)
49         list_splice(promote + i, &n->partial);
50
51     spin_unlock_irqrestore(&n->list_lock, flags);
52
53     /* Release empty slabs */
54     list_for_each_entry_safe(page, t, &discard, lru)
55         discard_slab(s, page);
56
57     if (slabs_node(s, node))
58         ret = 1;
59 }
60 return ret;
61 }

```

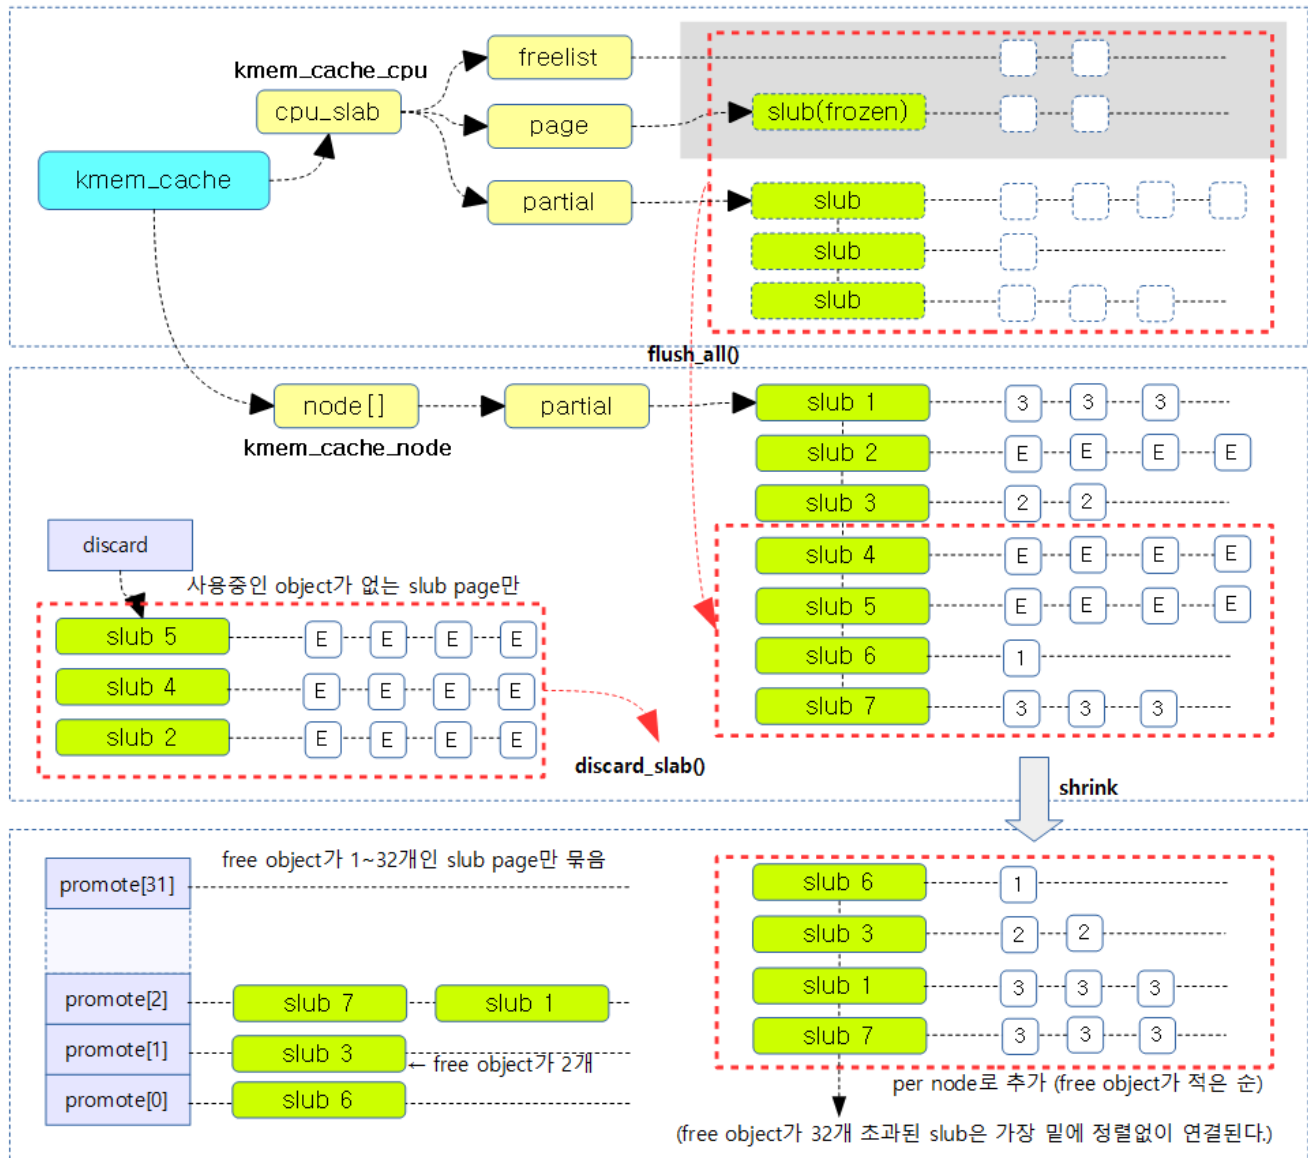
Perform the following procedure to find revocable slab pages in the requested slab cache and delete all of them.

- If none of the slab pages in the n->partial list have an object in use, the slab pages are returned to the buddy system.
- Fewer than 32 free objects in the slab pages of an n->partial list are ascending sorted.

- In line 13 of code, move all the per cpu slab pages in the slab cache to an n->partial list.
- In lines 14~17 of the code, traverse the nodes in the slab cache and initialize the discard list, and initialize the promote[] list by SHRINK_PROMOTE_MAX(32).
- In code lines 27~41, traverse the slab pages of the n->partial list, adding the slab pages with all the objects of free objects to the discard list, and if the number of free objects is less than SHRINK_PROMOTE_MAX(32), add them to the promote[free-1] list.
- In lines 47~48 of code, add the promote[i] list to the beginning of the n->partial list.
 - In the end, the free objects are sorted in the order of 1~32 and added to the lead.
 - Place the slab page with the fewest free objects at the top so that the slab page can be quickly removed from the list management when the object is assigned.
- In code lines 53~54, unload all the slab pages in the discard list and return them to the buddy system.
- In line 56~57 of the code, if there is a slab left on the node during slub debugging, the function will return 1 at the exit of the function.

The following illustration shows all the slab pages in the current slab cache being moved to an n->partial list and organized.

`_kmem_cache_shrink()` - per cpu 캐시의 모든 slab page를 per 노드로 옮기고 정리한다.



(http://jake.dothome.co.kr/wp-content/uploads/2016/06/kmem_cache_shrink-1.png)

kick_all_cpus_sync()

kernel/smp.c

```
01  /*  
02  * kick_all_cpus_sync - Force all cpus out of idle  
03  *  
04  * Used to synchronize the update of pm_idle function pointer. It's  
05  * called after the pointer is updated and returns after the dummy  
06  * callback function has been executed on all cpus. The execution of  
07  * the function can only happen on the remote cpus after they have  
08  * left the idle function which had been called via pm_idle function  
09  * pointer. So it's guaranteed that nothing uses the previous pointer  
10  * anymore.  
11  */  
  
1  void kick_all_cpus_sync(void)  
2  {  
3      /* Make sure the change is visible before we kick the cpus */  
4      smp_mb();  
5      smp_call_function(do_nothing, NULL, 1);  
6  }  
7  EXPORT_SYMBOL_GPL(kick_all_cpus_sync);
```

Use an IPI call to call each cpu to perform a dummy callback.

- pm_idle used to synchronize the update of the function pointer.
- do_nothing() is an empty function.

consultation

- Slab Memory Allocator -1- (Structure) (<http://jake.dothome.co.kr/slub/>) | Qc
- Slab Memory Allocator -2- (Initialize Cache) (http://jake.dothome.co.kr/kmem_cache_init) | Qc
- Slub Memory Allocator -3- (Create Cache) (<http://jake.dothome.co.kr/slub-cache-create>) | Qc
- Slub Memory Allocator -4- (Calculate Order) (<http://jake.dothome.co.kr/slub-order>) | Qc
- Slub Memory Allocator -5- | (<http://jake.dothome.co.kr/slub-slub-alloc>) Qc
- Slub Memory Allocator -6- (Assign Object) (<http://jake.dothome.co.kr/slub-object-alloc>) | Qc
- Slub Memory Allocator -7- (Object Unlocked) (<http://jake.dothome.co.kr/slub-object-free>) | Qc
- Slub Memory Allocator -8- (Drain/Flash Cache) (<http://jake.dothome.co.kr/slub-drain-flush-cache>) | Qc
- Slub Memory Allocator -9- (Cache Shrink) (<http://jake.dothome.co.kr/slub-cache-shrink>) | Sentence C – Current post
- Slub Memory Allocator -10- | (<http://jake.dothome.co.kr/slub-slub-free>) Qc
- Slub Memory Allocator -11- (Clear Cache (<http://jake.dothome.co.kr/slub-cache-destroy>)) | Qc
- Slub Memory Allocator -12- (Debugging Slub) (<http://jake.dothome.co.kr/slub-debug>) | Qc
- Slub Memory Allocator -13- (slabinfo) (<http://jake.dothome.co.kr/slub-slabinfo>) | 문c

LEAVE A COMMENT

Your email will not be published. Required fields are marked with *

Comments

name *

email *

Website

WRITE A COMMENT

◀ Slub Memory Allocator -11- (Clear Cache) (<http://jake.dothome.co.kr/slub-cache-destroy/>)

ZONE Type ▶ (<http://jake.dothome.co.kr/zone-types/>)

Munc Blog (2015 ~ 2024)