

Slub Memory Allocator -10-

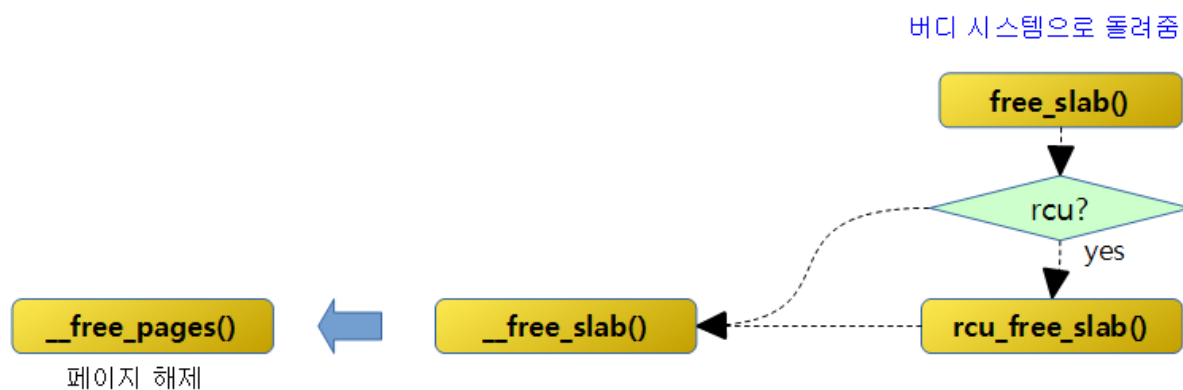
📅 2016-06-07 (<http://jake.dothome.co.kr/slub-slub-free/>) 👤 Moon Young-il

(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

<kernel v5.0>

Release slap pages

The following diagram shows the flow of the slub page to be released from the free_slab() function to return to the buddy system.



(http://jake.dothome.co.kr/wp-content/uploads/2016/06/free_slab-1.png)

free_slab()

mm/slub.c

```

1 | static void free_slab(struct kmem_cache *s, struct page *page)
2 | {
3 |     if (unlikely(s->flags & SLAB_TYPESAFE_BY_RCU)) {
4 |         call_rcu(&page->rcu_head, rcu_free_slab);
5 |     } else
6 |         __free_slab(s, page);
7 | }

```

Dismiss the slap page and return it to the buddy system.

- In line 3~4 of code, there is a small chance that lock-less RCU will be used to release the slap and reclaim it to the buddy system.
- In line 5~6 of the code, immediately reclaim the slap page to the buddy system.

rcu_free_slab()

mm/slub.c

```

1 | static void rcu_free_slab(struct rcu_head *h)
2 | {

```

```

3 | struct page *page = container_of(h, struct page, rcu_head);
4 |
5 | __free_slab(page->slab_cache, page);
6 | }

```

Reclaim the slab page to the buddy system. This function is called through the RCU.

__free_slab()

mm/slub.c

```

01 | static void __free_slab(struct kmem_cache *s, struct page *page)
02 | {
03 |     int order = compound_order(page);
04 |     int pages = 1 << order;
05 |
06 |     if (s->flags & SLAB_CONSISTENCY_CHECKS) {
07 |         void *p;
08 |
09 |         slab_pad_check(s, page);
10 |         for_each_object(p, s, page_address(page),
11 |                         page->objects)
12 |             check_object(s, page, p, SLUB_RED_INACTIVE);
13 |     }
14 |
15 |     mod_lruvec_page_state(page,
16 |         (s->flags & SLAB_RECLAIM_ACCOUNT) ?
17 |         NR_SLAB_RECLAIMABLE : NR_SLAB_UNRECLAIMABLE,
18 |         -pages);
19 |
20 |     __ClearPageSlabPfmemalloc(page);
21 |     __ClearPageSlab(page);
22 |
23 |     page->mapping = NULL;
24 |     if (current->reclaim_state)
25 |         current->reclaim_state->reclaimed_slab += pages;
26 |     memcg_uncharge_slab(page, order, s);
27 |     __free_pages(page, order);
28 | }

```

Dismiss the slab page and retrieve it to the buddy system.

- In line 3~4 of code, the slab page can be configured as a compound page. Therefore, we know the corresponding order value and the number of pages.
- If you used the SLAB_CONSISTENCY_CHECKS flag in lines 6~13 of the code, perform the consistency check on the slab page as follows:
 - Check the padding for the slab page.
 - Check the red-zone of each slab object to see if it is SLUB_RED_INACTIVE (0xbb).
- In code lines 15~18, decrement the NR_SLAB_RECLAIMABLE counter if it is a reclaimable slab cache, and the NR_SLAB_UNRECLAIMABLE counter by the number of pages if it is not.
- Clear PG_active and PG_slab on lines 20~23 of the code and unlink the page mapping.
- If the reclamation status of the current task is set in code lines 24~25, add the reclaimed_slab by the number of pages.
- In line 26 of the code, the memory cgroup tells us that the slab page has been recalled, so that it is reduced in counting.
- At line 27 of code, the buddy system retrieves the number of pages in order.

memcg_uncharge_slab()

mm/slab.h

```

1 | static __always_inline void memcg_uncharge_slab(struct page *page, int order,
2 |                                                 struct kmem_cache *s)
3 | {
4 |     if (!memcg_kmem_enabled())
5 |         return;
6 |     memcg_kmem_uncharge(page, order);
7 | }

```

If memcg_kmem is enabled while using the CONFIG_MEMCG_KMEM kernel options, the memory cgroup will notify you that the slab page has been recalled and reduce it from counting. However, if it is a root cache, it will not be processed and will exit the routine.

consultation

- Slab Memory Allocator -1- (Structure) (<http://jake.dothome.co.kr/slub/>) | Qc
- Slab Memory Allocator -2- (Initialize Cache) (http://jake.dothome.co.kr/kmem_cache_init) | Qc
- Slub Memory Allocator -3- (Create Cache) (<http://jake.dothome.co.kr/slub-cache-create>) | Qc
- Slub Memory Allocator -4- (Calculate Order) (<http://jake.dothome.co.kr/slub-order>) | Qc
- Slub Memory Allocator -5- | (<http://jake.dothome.co.kr/slub-slub-alloc>) Qc
- Slub Memory Allocator -6- (Assign Object) (<http://jake.dothome.co.kr/slub-object-alloc>) | Qc
- Slub Memory Allocator -7- (Object Unlocked) (<http://jake.dothome.co.kr/slub-object-free>) | Qc
- Slub Memory Allocator -8- (Drain/Flash Cache) (<http://jake.dothome.co.kr/slub-drain-flush-cache>) | Qc
- Slub Memory Allocator -9- (Cache Shrink) (<http://jake.dothome.co.kr/slub-cache-shrink>) | Qc
- Slub Memory Allocator -10- | (<http://jake.dothome.co.kr/slub-slub-free>) Sentence C – Current post
- Slub Memory Allocator -11- (Clear Cache (<http://jake.dothome.co.kr/slub-cache-destroy>)) | Qc
- Slub Memory Allocator -12- (Debugging Slub) (<http://jake.dothome.co.kr/slub-debug>) | Qc
- Slub Memory Allocator -13- (slabinfo) (<http://jake.dothome.co.kr/slub-slabinfo>) | 문c

LEAVE A COMMENT

Your email will not be published. Required fields are marked with *

Comments

name *

email *

Website

WRITE A COMMENT

◀ Slub Memory Allocator -6- (Assign Object) (<http://jake.dothome.co.kr/slub-object-alloc/>)

Slub Memory Allocator -7- (Object Unlocked) ▶ (<http://jake.dothome.co.kr/slub-object-free/>)

Munc Blog (2015 ~ 2024)