

# Swap Entries

📅 2017-01-10 (<http://jake.dothome.co.kr/swap-entry/>) 👤 Moon Young-il  
(<http://jake.dothome.co.kr/author/admin/>) 📁 Linux Kernel (<http://jake.dothome.co.kr/category/linux/>)

<kernel v5.0>

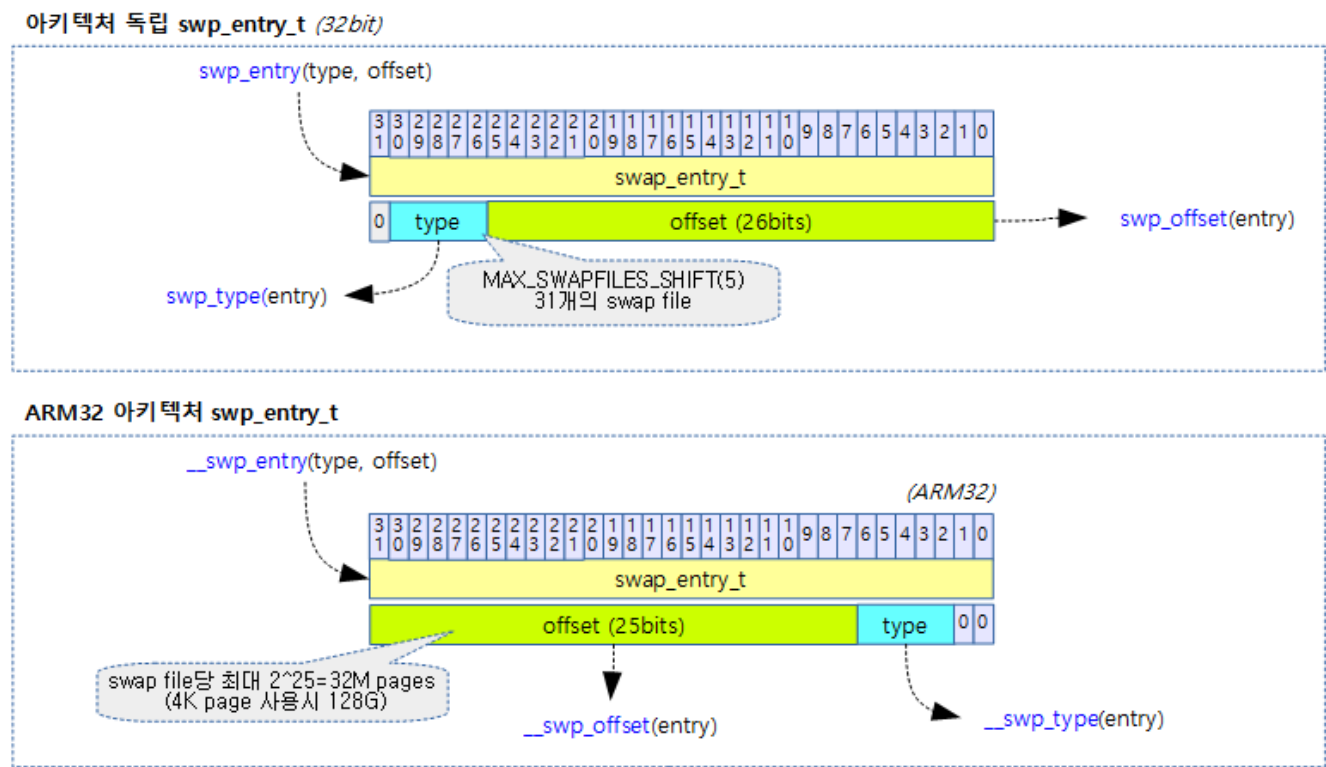
## Swap Entries

In kernel v4.20-rc1, the management of swap entries in the Linux kernel has been changed from using Radix Tree Exceptional to using XArray.

- See: xarray: Replace exceptional entries  
(<https://github.com/torvalds/linux/commit/3159f943aafdbacb2f94c38fdaadabf2bbde2a14>)

### swap\_entry\_t Structure

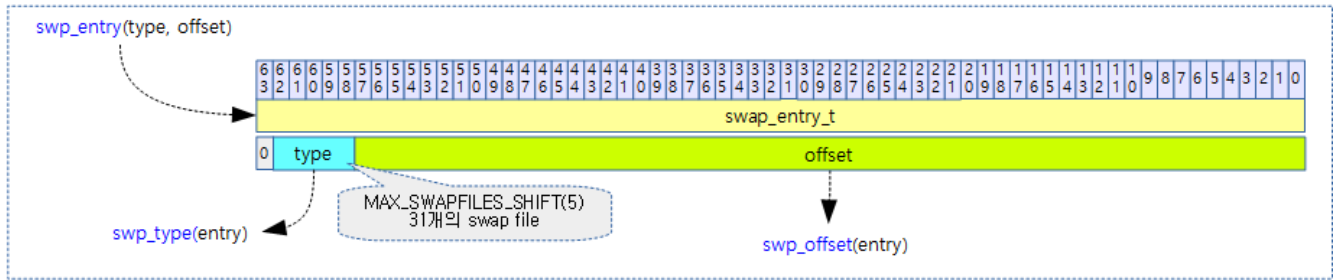
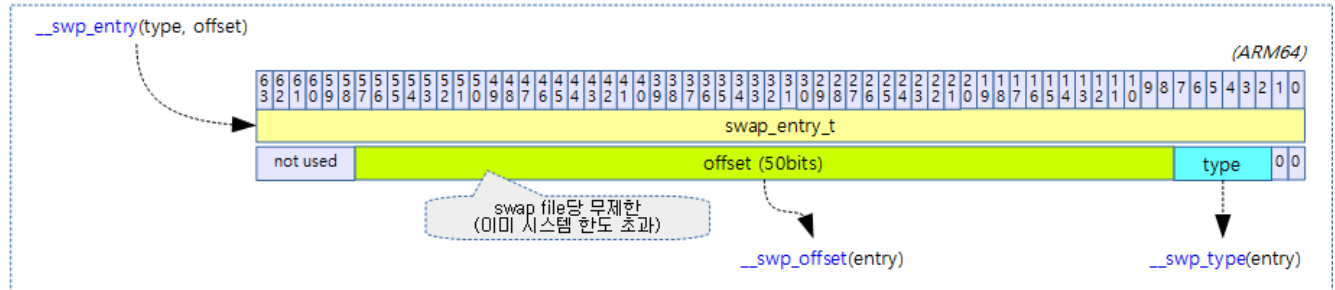
The following figure shows the swap entry structure running on a 32-bit system.



([http://jake.dothome.co.kr/wp-content/uploads/2017/01/swp\\_entry-1c.png](http://jake.dothome.co.kr/wp-content/uploads/2017/01/swp_entry-1c.png))

- This value is stored in the ARM Linux PTE entry, not the ARM H/W PTE entry, which means that this page is a swap entry.

The following figure shows the swap entry structure running on a 64-bit system.

아키텍처 독립 `swp_entry_t` (64bit)ARM64 아키텍처 `swp_entry_t`

(http://jake.dothome.co.kr/wp-content/uploads/2017/01/swp\_entry-2.png)

## swp\_entry\_t Structure

include/linux/mm\_types.h

```

1  /*
2  * A swap entry has to fit into a "unsigned long", as the entry is hidden
3  * in the "index" field of the swapper address space.
4  */

1  typedef struct {
2      unsigned long val;
3  } swp_entry_t;

```

## Architecture-independent swap entry

### swp\_entry()

include/linux/swapops.h

```

01  /*
02  * Store a type+offset into a swp_entry_t in an arch-independent format
03  */
04  static inline swp_entry_t swp_entry(unsigned long type, pgoff_t offset)
05  {
06      swp_entry_t ret;
07
08      ret.val = (type << SWP_TYPE_SHIFT(ret)) | (offset & SWP_OFFSET_MASK);
09      return ret;
10  }

```

Configure the swap entry with the offset value and the type value.

### swp\_type()

include/linux/swapops.h

```

1  /*
2  * Extract the `type' field from a swp_entry_t. The swp_entry_t is in
3  * arch-independent format
4  */

1  static inline unsigned swp_type(swp_entry_t entry)
2  {
3      return (entry.val >> SWP_TYPE_SHIFT);
4  }
```

Returns a type value from the swap entry.

## swp\_offset()

include/linux/swapops.h

```

1  /*
2  * Extract the `offset' field from a swp_entry_t. The swp_entry_t is in
3  * arch-independent format
4  */

1  static inline pgoff_t swp_offset(swp_entry_t entry)
2  {
3      return entry.val & SWP_OFFSET_MASK;
4  }
```

Returns the value of offset from the swap entry.

## SWP\_TYPE\_SHIFT() & SWP\_OFFSET\_MASK()

include/linux/swapops.h

```

01  /*
02  * swapcache pages are stored in the swapper_space radix tree. We want
03  * to get good packing density in that tree, so the index should be dense i
04  * n the low-order bits.
05  *
06  * We arrange the `type' and `offset' fields so that `type' is at the se
07  * ven high-order bits of the swp_entry_t and `offset' is right-aligned in t
08  * he remaining bits. Although `type' itself needs only five bits, we allo
09  * w for shmem/tmpfs to shift it all up a further two bits: see swp_to_radix_e
10  * ntry().
11  * swp_entry_t's are *never* stored anywhere in their arch-dependent for
12  * mat.
13  */

1  #define SWP_TYPE_SHIFT      (BITS_PER_XA_VALUE - MAX_SWAPFILES_SHIFT)
2  #define SWP_OFFSET_MASK    ((1UL << SWP_TYPE_SHIFT) - 1)
```

## Macro Constants

include/linux/swap.h

```

1  /*
2  * MAX_SWAPFILES defines the maximum number of swaptypes: things which c
   an
3  * be swapped to. The swap type and the offset into that swap type are
4  * encoded into pte's and into pgoff_t's in the swapcache. Using five b
   its
5  * for the type means that the maximum number of swapcache pages is 27 b
   its
6  * on 32-bit-pgoff_t architectures. And that assumes that the architect
   ure packs
7  * the type/offset into the pte as 5/27 as well.
8  */

1  #define MAX_SWAPFILES_SHIFT      5

```

include/linux/swap.h

```

1  /*
2  * NUMA node memory migration support
3  */

1  #ifdef CONFIG_MIGRATION
2  #define SWP_MIGRATION_NUM 2
3  #define SWP_MIGRATION_READ      (MAX_SWAPFILES + SWP_HWPOISON_NUM)
4  #define SWP_MIGRATION_WRITE    (MAX_SWAPFILES + SWP_HWPOISON_NUM + 1)
5  #else
6  #define SWP_MIGRATION_NUM 0
7  #endif

```

## ARM32 Swap Entry

### \_\_swp\_entry()

arch/arm/include/asm/pgtable.h

```

1  #define __swp_entry(type, offset) ((swp_entry_t) { ((type) << __SWP_TYPE_
   SHIFT) | ((offset) << __SWP_OFFSET_SHIFT) })

```

The offset value and the type value are used to construct the arm swap entry.

### \_\_swp\_type()

arch/arm/include/asm/pgtable.h

```

1  #define __swp_type(x)      (((x).val >> __SWP_TYPE_SHIFT) & __SWP_T
   YPE_MASK)

```

Arm swap entry.

### \_\_swp\_offset()

arch/arm/include/asm/pgtable.h

```

1  #define __swp_offset(x)      ((x).val >> __SWP_OFFSET_SHIFT)

```

Arm Swap returns the offset value from the entry.

## Macro Constants

arch/arm/include/asm/pgtable.h

```

01  /*
02  * Encode and decode a swap entry.  Swap entries are stored in the Linux
03  * page tables as follows:
04  *
05  *   3 3 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1 1
06  *   1 0 9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0
07  *   <----- offset -----> < type -> 0 0
08  *
09  * This gives us up to 31 swap files and 128GB per swap file.  Note that
10  * the offset field is always non-zero.
11  */

1  #define __SWP_TYPE_SHIFT          2
2  #define __SWP_TYPE_BITS          5
3  #define __SWP_TYPE_MASK          ((1 << __SWP_TYPE_BITS) - 1)
4  #define __SWP_OFFSET_SHIFT       (__SWP_TYPE_BITS + __SWP_TYPE_SHIFT)

```

## ARM64 Swap Entry

### \_\_swp\_entry()

arch/arm64/include/asm/pgtable.h

```

1  #define __swp_entry(type,offset) ((swp_entry_t) { ((type) << __SWP_TYPE_
    SHIFT) | ((offset) << __SWP_OFFSET_SHIFT) })

```

The offset value and the type value make up the arm64 swap entry.

### \_\_swp\_type()

arch/arm64/include/asm/pgtable.h

```

1  #define __swp_type(x)              (((x).val >> __SWP_TYPE_SHIFT) & __SWP_T
    YPE_MASK)

```

arm64 swap entry.

### \_\_swp\_offset()

arch/arm64/include/asm/pgtable.h

```

1  #define __swp_offset(x)            (((x).val >> __SWP_OFFSET_SHIFT) & __SWP
    _OFFSET_MASK)

```

arm64 swap entry returns a value of 50 bit offset.

## Macro Constants

arch/arm64/include/asm/pgtable.h

```

1  /*
2  *  Encode and decode a swap entry:
3  *      bits 0-1:      present (must be zero)
4  *      bits 2-7:      swap type
5  *      bits 8-57:     swap offset
6  *      bit 58:        PTE_PROT_NONE (must be zero)
7  */

1  #define __SWP_TYPE_SHIFT      2
2  #define __SWP_TYPE_BITS      6
3  #define __SWP_OFFSET_BITS    50
4  #define __SWP_TYPE_MASK      ((1 << __SWP_TYPE_BITS) - 1)
5  #define __SWP_OFFSET_SHIFT   (__SWP_TYPE_BITS + __SWP_TYPE_SHIFT)
6  #define __SWP_OFFSET_MASK    ((1UL << __SWP_OFFSET_BITS) - 1)

1  |

```

## Swap PTE Entry Identification

### is\_swap\_pte()

include/linux/swapops.h

```

1  /* check whether a pte points to a swap entry */
2  static inline int is_swap_pte(pte_t pte)
3  {
4      return !pte_none(pte) && !pte_present(pte);
5  }

```

Returns whether it is a swapped pte entry.

- If the PTE is not NONE and there is no PRESENT setting, it is in a swap state.

## Swap Cache

Shared pages with slots reserved for backing storage are considered swap caches. The swap cache differs from the file cache in the following two ways.

- page->mapping uses &swapper\_space[]. (address\_space)
- Use the add\_to\_swap\_cache() function instead of add\_to\_page\_cache().

---

### LEAVE A COMMENT

Your email will not be published. Required fields are marked with \*

Comments

name \*

email \*

Website

WRITE A COMMENT

◀ Exception -4- (ARM32 VFP & FPE) (<http://jake.dothome.co.kr/vfp-fpe/>)

Exception -6- (MM Fault Handler) ▶ (<http://jake.dothome.co.kr/mm-fault/>)

Munc Blog (2015 ~ 2024)